

# Data Science & ML Course

## Lesson #4 [Part #2] Introduction to Pandas

Ivanovitch Silva  
October, 2018



# Agenda

---

- About the two core pandas types: **dataframes** and **series**
- How to **select** data using row and column labels
- A variety of methods for **exploring data** with pandas
- How to **assign** data using various techniques in pandas
- How to use **boolean indexing** with pandas for selection and assignment

# Update from repository

---

```
git clone https://github.com/ivanovitchm/datascience2machinelearning.git
```

Or ....

```
git pull
```



# Understanding Pandas & Numpy

---

- Numpy
  - a. Lack support for column names
  - b. Support for only one data type per ndarray
  - c. There are lots of low level methods, however there are many common analysis patterns that don't have pre-built methods.

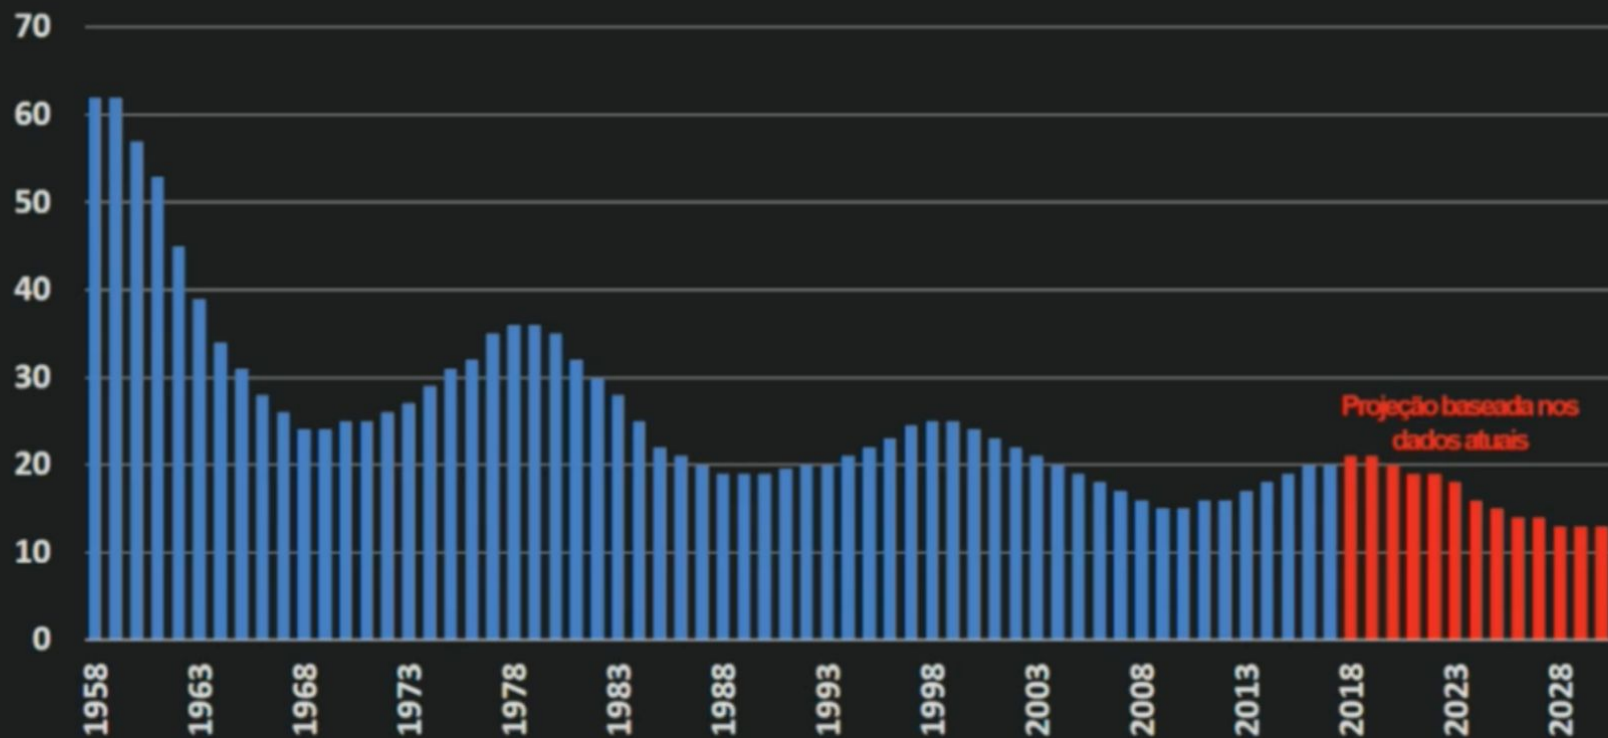
**The pandas library** provides solutions to all of these pain points and more. Pandas is not so much a replacement for NumPy as an **extension of NumPy**.

THE WORLD'S LARGEST CORPORATIONS  
**FORTUNE**

GLOBAL



## Tempo médio de permanência de uma empresa no S&P 500 (em anos)



Fonte: INNOSIGHT, Richard N. Foster, Standard & Poor's

# UMA EMPRESA DO S&P 500 ESTÁ SENDO SUBSTITUÍDA A CADA DUAS SEMANAS

Richard Foster

# The dataset

	rank	revenues	revenue_change	profits	assets	profit_change	ceo	industry	sector	previous_rank
<b>Walmart</b>	1	485873	0.8	13643.0	198825	-7.2	C. Douglas McMillon	General Merchandisers	Retailing	1
<b>State Grid</b>	2	315199	-4.4	9571.3	489838	-6.2	Kou Wei	Utilities	Energy	2
<b>Sinopec Group</b>	3	267518	-9.1	1257.9	310726	-65.0	Wang Yupu	Petroleum Refining	Energy	4
<b>China National Petroleum</b>	4	262573	-12.3	1867.5	585619	-73.7	Zhang Jianhua	Petroleum Refining	Energy	3
<b>Toyota Motor</b>	5	254694	7.7	16899.3	437575	-12.3	Akio Toyoda	Motor Vehicles and Parts	Motor Vehicles & Parts	8

```
import pandas as pd
f500 = pd.read_csv("f500.csv", index_col=0)
f500.index.name = None
```



# Introducing Dataframes

The diagram illustrates a DataFrame structure with the following components:

- Column Labels:** rank, revenues, profits, country
- Index Axis:** Indicated by a red arrow pointing downwards.
- Row Labels:** Walmart, State Grid, Sinopec Group, China Natural Petroleum, Toyota Motor
- Column Axis:** Indicated by a red arrow pointing to the right.
- Data Types:**
  - Integer Type:** Indicated by green arrows pointing to the 'rank' and 'revenues' columns.
  - Float Type:** Indicated by a green arrow pointing to the 'profits' column.
  - String Type:** Indicated by green arrows pointing to the 'country' and 'Row Labels'.

	rank	revenues	profits	country
Walmart	1	485873	13643.0	USA
State Grid	2	315199	9571.3	China
Sinopec Group	3	267518	1257.9	China
China Natural Petroleum	4	262573	1867.5	China
Toyota Motor	5	254694	16899.3	Japan

# Introducing Dataframes

---

```
# put your code here  
f500.info(,)
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 500 entries, Walmart to AutoNation  
Data columns (total 16 columns):  
rank                500 non-null int64  
revenues            500 non-null int64  
revenue_change      498 non-null float64  
profits            499 non-null float64  
assets             500 non-null int64  
profit_change      436 non-null float64  
ceo                500 non-null object  
industry           500 non-null object  
sector            500 non-null object  
previous_rank      500 non-null int64  
country           500 non-null object  
hq_location        500 non-null object  
website            500 non-null object  
years_on_global_500_list 500 non-null int64  
employees          500 non-null int64  
total_stockholder_equity 500 non-null int64  
dtypes: float64(3), int64(7), object(6)  
memory usage: 66.4+ KB
```

f500.head()  
f500.tail()

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
f500_selection	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

```
f500_selection.loc[:, "rank"]
```

Walmart	1
State Grid	2
Sinopec Group	3
China Natural Petroleum	4
Toyota Motor	5

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

		country	rank
f500_selection.loc[:, ["country", "rank"]]	Walmart	USA	1
	State Grid	China	2
	Sinopec Group	China	3
	China Natural Petroleum	China	4
	Toyota Motor	Japan	5

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

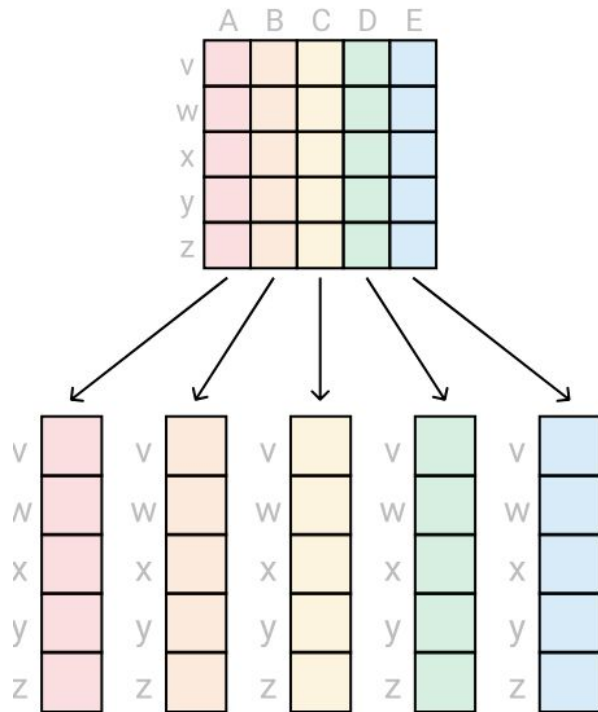
		rank	revenues	profits
f500_selection.loc[:, "rank": "profits"]	Walmart	1	485873	13643.0
	State Grid	2	315199	9571.3
	Sinopec Group	3	267518	1257.9
	China Natural Petroleum	4	262573	1867.5
	Toyota Motor	5	254694	16899.3

# Column selection shortcuts

---

Select by Label	Explicit Syntax	Common Shorthand	Other Shorthand
Single column	<code>df.loc[:, "col1"]</code>	<code>df["col1"]</code>	<code>df.col1</code>
List of columns	<code>df.loc[:, ["col1", "col7"]]</code>	<code>df[["col1", "col7"]]</code>	
Slice of columns	<code>df.loc[:, "col1":"col4"]</code>		

# Selecting Items from a Series by Label



Original  
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

## Code

```
single_col = df["D"]
```

## Result

V	
W	
X	
Y	
Z	

**single\_col** is a  
series object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
single_row = df.head(1)
```

A	
B	
C	
D	
E	

**single\_row** is a  
series object



Original  
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

## Code

```
multi_cols = df[["A", "C", "D"]]
```

## Result

	A	C	D
V			
W			
X			
Y			
Z			

**multi\_cols** is a  
dataframe object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
multi_rows = df.head(3)
```

	A	B	C	D	E
V					
W					
X					

**multi\_rows** is a  
dataframe object

# Dataframe vs Series

---

	<b>Series</b>	<b>DataFrame</b>
<b>Dimensions</b>	One	Two
<b>Has 'index' axis</b>	Yes	Yes
<b>Has 'columns' axis</b>	No	Yes
<b>Number of dtypes</b>	One	Many (one per column)

# Series and Dataframe Describe Methods

---

```
revs = f500["revenues"]  
print(revs.describe())
```


```
count      500.000000  
mean      55416.358000  
std       45725.478963  
min       21609.000000  
25%       29003.000000  
50%       40236.000000  
75%       63926.750000  
max       485873.000000  
Name: revenues, dtype: float64
```

```
print(f500["assets"].describe())
```

```
count      5.000000e+02  
mean      2.436323e+05  
std       4.851937e+05  
min       3.717000e+03  
25%      3.658850e+04  
50%      7.326150e+04  
75%      1.805640e+05  
max      3.473238e+06  
Name: assets, dtype: float64
```

`DataFrame.method(axis=0)`  
or  
`DataFrame.method(axis="index")`

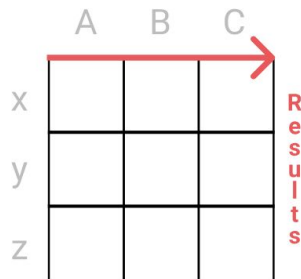
Calculates along the **row** axis



Calculates result for  
each **column**.

`DataFrame.method(axis=1)`  
or  
`DataFrame.method(axis="column")`

Calculates along the **column** axis



Calculates result for  
each **row**.

## More data exploration methods

```
medians = f500[["revenues", "profits"]].median(axis=0)
# we could also use .median(axis="index")
print(medians)
```

```
revenues    40236.0
profits      1761.6
dtype: float64
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	485873
State Grid	2	315199
Sinopec Group	3	267518
China National Petroleum	4	262573
Toyota Motor	5	254694

```
>>> top5_rank_revenue["revenues"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	0
China National Petroleum	4	0
Toyota Motor	5	0

## Assignment with Pandas

# Assignment with Pandas

---

```
>>> top5_rank_revenue.loc["Sinopec Group", "revenues"] = 999
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	999
China National Petroleum	4	0
Toyota Motor	5	0

# Add a new column

---

```
>>> top5_rank_revenue["year_founded"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0

# Add a new row

---

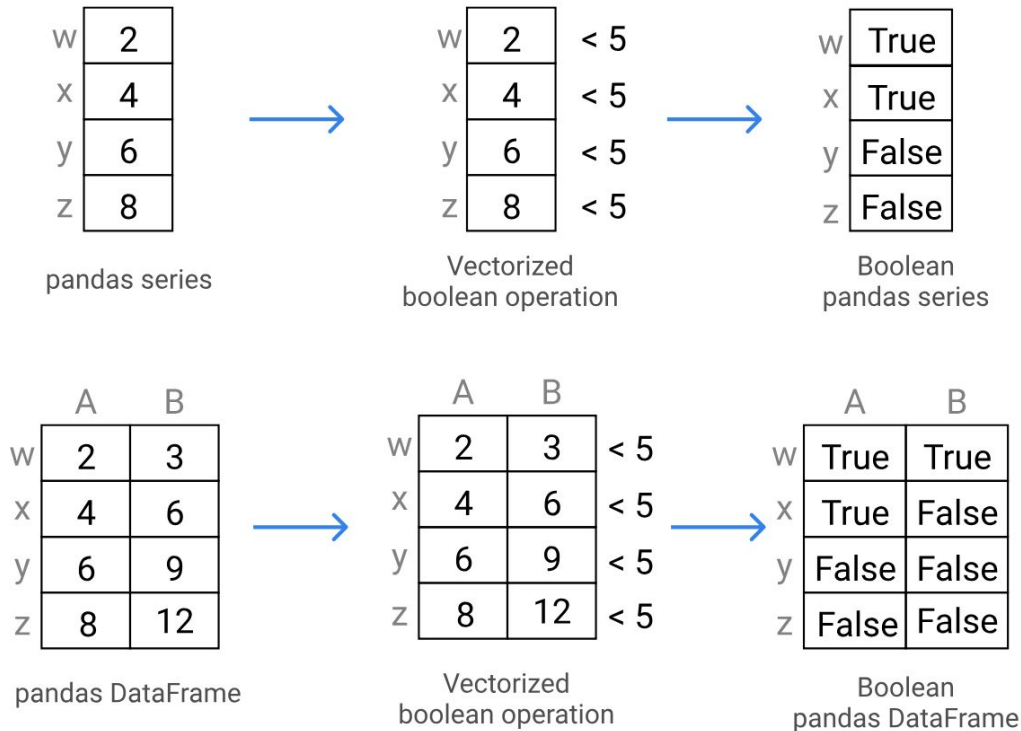
```
>>> top5_rank_revenue.loc["My New Company"] = 555
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0
My New Company	555	555	555



# Using boolean indexing with pandas objects



# Using boolean indexing with pandas objects

```
result = df.loc[num_bool, "name"]
```

		name		
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↑
				<b>result</b>
				Rahul
				Sarah

```
result = df[num_bool]
```

		name	num	
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↑
				<b>result</b>
				name num
				Rahul 8
				Sarah 8

# Using boolean arrays to assign values

---

```
f500.loc[f500["sector"] == "Motor Vehicles & Parts", "sector"] = "Motor Vehicles and Parts"
```

# Challenge

---

## Finding top performers by country

```
>>> top_3_countries = f500["country"].value_counts().head(3)
```

```
>>> print(top_3_countries)
```

```
USA      132
```

```
China    109
```

```
Japan     51
```

```
Name: country, dtype: int64
```

# Lesson\_04 Part #2

## Section 1.12



# Exploring Data With Pandas

# Agenda

---

- Select columns, rows and individual items using their integer location.
- Work with integer axis labels.
- How to use pandas methods to produce boolean arrays.
- Use boolean operators to combine boolean comparisons to perform more complex analysis.
- Use index labels to align data.
- Use aggregation to perform advanced analysis using loops.

# Introduction (Pandas vs Numpy)


```
ndarray[2,0]
```

*located at row 2, column 0*


```
ndarray[1]
```

*located at row 1*

Numpy

	A	B	C
x			
y			
z			

```
df.loc["z","A"]
```

*located at row with label z,  
column with label A*

	A	B	C
x			
y			
z			

```
df.loc["y"]
```

*located at row with label y*

Pandas



# Using iloc to select by integer position

	A	B	C
x			
y			
z			

```
df.iloc[2,0]
```

	A	B	C
x			
y			
z			

```
df.iloc[1]
```

```
first_column = f500.iloc[:,0]
print(first_column)
```

```
0           Walmart
1       State Grid
2       Sinopec Group
...
497  Wm. Morrison Supermarkets
498           TUI
499       AutoNation
Name: company, dtype: object
```

# Slicing with iloc

With `loc[]`, the ending slice is **included**.

With `iloc[]`, the ending slice is **not included**.

```
1 | f500[1:4].
```

	rank	revenues
State Grid	2	315199
Sinopec Group	3	267518
China National Petroleum	4	262573

```
1 | f500.iloc[1:4].
```

	rank	revenues
State Grid	2	315199
Sinopec Group	3	267518
China National Petroleum	4	262573

# Loc vs iLoc

`df.iloc[1]`

`iloc[1]` uses the integer position of the row to select the second row

	A	B	C
0			
1			
2			

`df.iloc[1]`

`iloc[1]` uses the integer position of the row to select the second row

	A	B	C
0			
2			
1			

`df.loc[1]`

`loc[1]` uses the label of the row to select the row with an axis label of 1.

	A	B	C
0			
1			
2			

`df.loc[1]`

`loc[1]` uses the label of the row to select the row with an axis label of 1.

	A	B	C
0			
2			
1			

# Using pandas methods to create boolean masks

---

```
>>> is_california = usa["hq_location"].str.endswith("CA")
```

```
>>> print(is_california.head())
```

```
0      False
```

```
7      False
```

```
8       True
```

```
9      False
```

```
10     True
```

```
Name: hq_location, dtype: bool
```

```
0      Bentonville, AR
```

```
7      Omaha, NE
```

```
8      Cupertino, CA
```

```
9      Irving, TX
```

```
10     San Francisco, CA
```

```
Name: hq_location, dtype: object
```

# Using boolean operators to select items

	company	revenues	country
0	Walmart	485873	USA
1	State Grid	315199	China
2	Sinopec Group	267518	China
3	China Nation...	262573	China
4	Toyota Motor	254694	Japan

f500\_sel

```
over_265 = f500_sel["revenues"] > 265000
china = f500_sel["country"] == "China"
```

0	True
1	True
2	True
3	False
4	False

over\_265

0	False
1	True
2	True
3	True
4	False

china

```
combined = over_265 & china
```

0	True	&	0	False	=	0	False
1	True	&	1	True	=	1	True
2	True	&	2	True	=	2	True
3	False	&	3	True	=	3	False
4	False	&	4	False	=	4	False

over\_265      china      combined

# Using boolean operators to select items

```
final_cols = ["company", "revenues"]  
result = f500_sel.loc[combined, final_cols]
```

		company	revenues	country		company	revenues
0	False	0	Walmart	485873	USA		
1	True →	1	State Grid	315199	China	→ 1	State Grid 315199
2	True →	2	Sinopec Group	267518	China	→ 2	Sinopec Group 267518
3	False	3	China Nation...	262573	China		
4	False	4	Toyota Motor	254694	Japan		

**combined** **f500\_sel** **result**

# Pandas Index Alignment

	fruit_veg	qty		
tomato	fruit	4	corn	yellow
carrot	veg	2	carrot	orange
lime	fruit	4	tomato	red
corn	veg	1	lime	green
eggplant	veg	2	eggplant	purple
food			colors	

	fruit_veg	qty	color
tomato	fruit	4	red
carrot	veg	2	orange
lime	fruit	4	green
corn	veg	1	yellow
eggplant	veg	2	purple
food			

```
food["color"] = colors
```

	fruit_veg	qty		
tomato	fruit	4	corn	yellow
carrot	veg	2	carrot	orange
lime	fruit	4	tomato	red
corn	veg	1	lime	green
eggplant	veg	2	eggplant	purple
food			colors	

# Pandas Index Alignment

---

arugula	rocket
eggplant	aubergine
corn	maize
	<b>alt_name</b>

`food["alt_name"] = alt_name`

	fruit_veg	qty	color	alt_name
tomato	fruit	4	red	NaN
carrot	veg	2	orange	NaN
lime	fruit	4	green	NaN
corn	veg	1	yellow	maize
eggplant	veg	2	purple	aubergine
	<b>food</b>			



# Using Loops in Pandas

---

```
>>> print(df)
```

	A	B	C
x	6	1	0
y	1	8	8
z	3	8	7

```
>>> for i in df:  
    print(i)
```

A  
B  
C

Because one of the key benefits of pandas is that it has vectorized methods to work with data more efficiently, **we want to avoid using loops wherever we can**

# Challenge: calculating return on assets by sector

---

```
{'Aerospace & Defense': 'Lockheed Martin',
'Apparel': 'Nike',
'Business Services': 'Adecco Group',
'Chemicals': 'LyondellBasell Industries',
'Energy': 'National Grid',
'Engineering & Construction': 'Pacific Construction Group',
'Financials': 'Berkshire Hathaway',
'Food & Drug Stores': 'Publix Super Markets',
'Food, Beverages & Tobacco': 'Philip Morris International',
'Health Care': 'Gilead Sciences',
'Hotels, Restaurants & Leisure': 'McDonald\xe2\x80\x99s',
'Household Products': 'Unilever',
'Industrials': '3M',
'Materials': 'CRH',
'Media': 'Disney',
'Motor Vehicles & Parts': 'Subaru',
'Retailing': 'H & M Hennes & Mauritz',
'Technology': 'Accenture',
'Telecommunications': 'KDDI',
'Transportation': 'Delta Air Lines',
'Wholesalers': 'McKesson'}
```

$$\text{return on assets} = \frac{\text{profits}}{\text{assets}}$$

# Lesson\_04 Part #2

## Section 2.10



# Challenge (...)

---

<https://github.com/torvalds/linux>

```
git log --encoding=latin-1 --pretty="%at,%aN" > log.csv
```

How many people contributed to the project?  
Top #10 contributors?

