

# Data Science & ML Course

## Lesson #15 Statistics Fundamentals II

Ivanovitch Silva  
November, 2018



# Agenda

---

- Frequency Distributions
  - Sorting frequency distribution tables
  - Percentiles and percentiles ranks
  - Information loss
- Visualizing Distributions
  - Bar, Pie, Histograms plots
  - Skewed distributions
  - Symmetrical Distributions
- Comparing Frequency Distribution

# Update from repository

---

```
git clone https://github.com/ivanovitchm/datascience2machinelearning.git
```

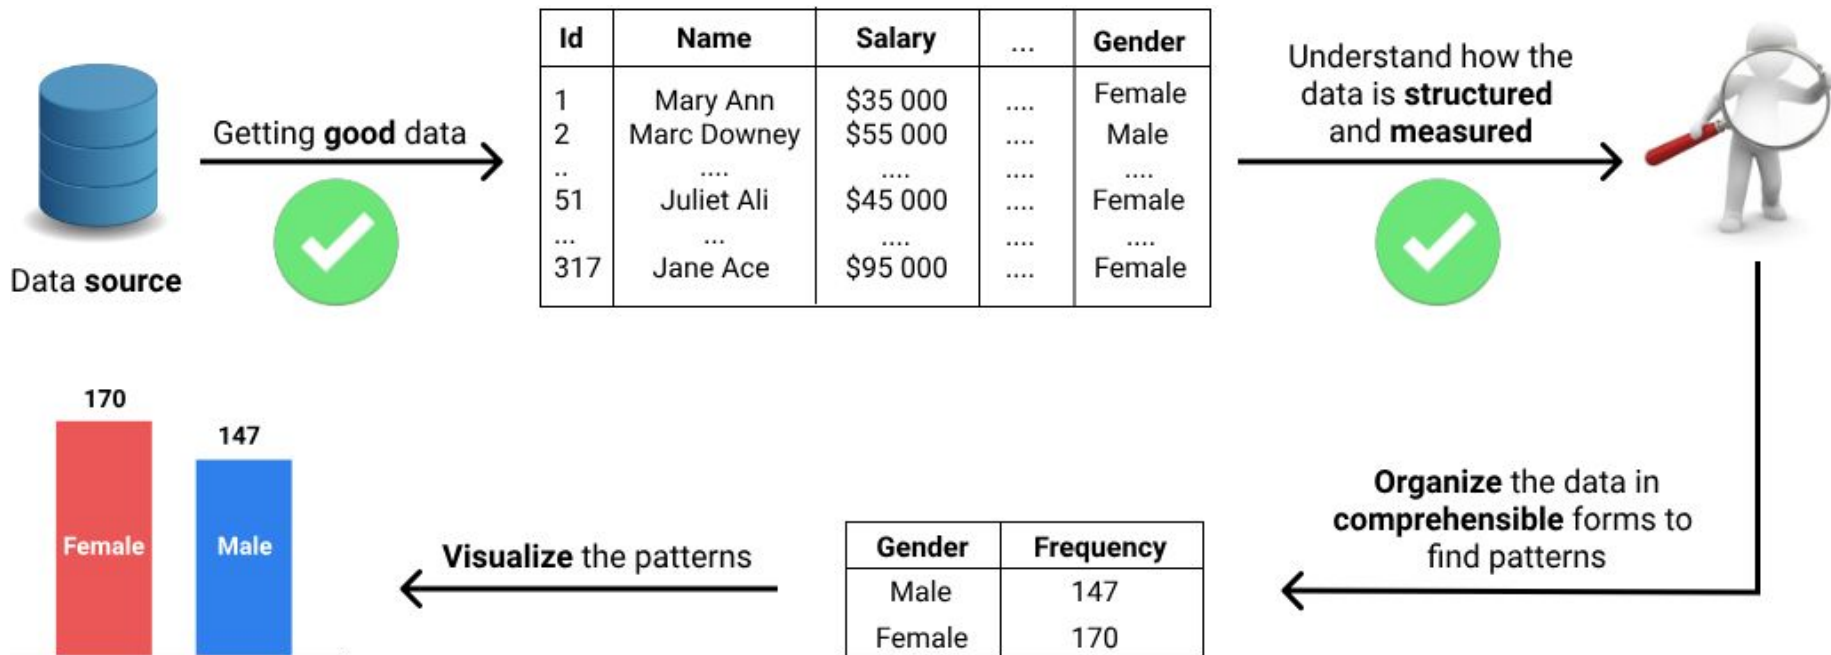
Or ....

```
git pull
```



# PREVIOUSLY ON...

(1) collecting data (2) understanding its structure and how it's measured

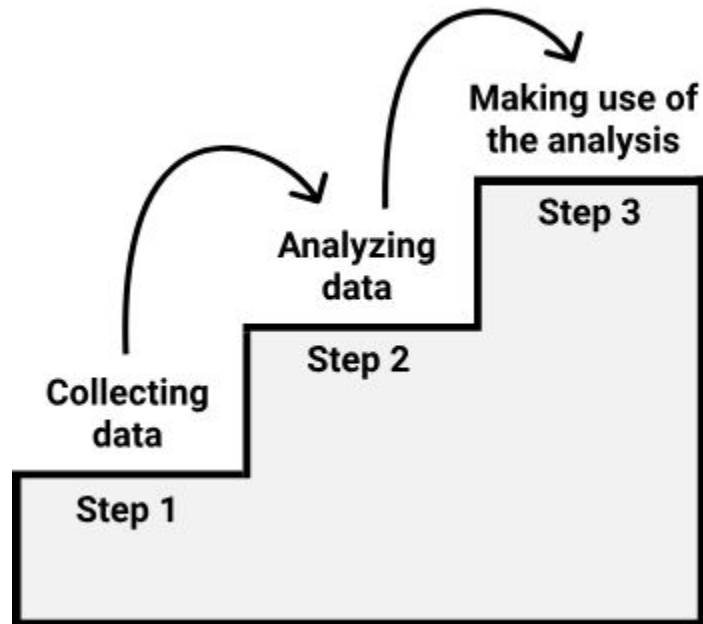


# Simplifying Data

---

We collect data to analyze it, and we analyze it for different purposes:

- To describe phenomena in the world (science).
- To make better decisions (industries).
- To improve systems (engineering).
- To describe different aspects of our society (data journalism); etc.



# Frequency Distribution Tables

Our capacity to understand a data set just by looking at it in a table format is limited

Position	Frequency
G	60
F	33
C	25
G/F	13
F/C	12

Unique values

Frequencies

_	Name	Team	Pos	Height	Weight	BMI	Birth_Place	Birthdate	Age	College	Experience	Games Played	MIN
0	Aerial Powers	DAL	F	183	71.0	21.200991	US	January 17, 1994	23	Michigan State	2	8	173
1	Alana Beard	LA	G/F	185	73.0	21.329438	US	May 14, 1982	35	Duke	12	30	947
2	Alex Bentley	CON	G	170	69.0	23.875433	US	October 27, 1990	26	Penn State	4	26	617
3	Alex Montgomery	SAN	G/F	185	84.0	24.543462	US	December 11, 1988	28	Georgia Tech	6	31	721
4	Alexis Jones	MIN	G	175	78.0	25.469388	US	August 5, 1994	23	Baylor	R	24	137

# Sorting Frequency Distribution Tables

```
1 wnba.Pos.value_counts(.)
```

G	60
F	33
C	25
G/F	13
F/C	12

Name: Pos, dtype: int64

Nominal Scale

```
1 wnba.Height.value_counts(.)
```

188	20
193	18
175	16
185	15
191	11
183	11
173	11
196	9
178	8
180	7
170	6
198	5
201	2
168	2
206	1
165	1

How many players are under 170 cm?

How many players are very tall (over 185)?

Are there any players below 160 cm?

Ratio Scale

We can tell the direction of difference

Name: Height, dtype: int64

# Sorting Tables for Ordinal Variables

---

The sorting techniques learned in the previous screen can't be used for ordinal scales where the measurement is done using words.

Condition	Label
points $\leq$ 20	very few points
20 < points $\leq$ 80	few points
80 < points $\leq$ 150	many points
points > 150	a lot of points

	Name	PTS	PTS_ordinal_scale
0	Aerial Powers	93	many points
1	Alana Beard	217	a lot of points
2	Alex Bentley	218	a lot of points
3	Alex Montgomery	188	a lot of points
4	Alexis Jones	50	few points



# Sorting Tables for Ordinal Variables

---

```
>> wnba['PTS_ordinal_scale'].value_counts()
```

```
a lot of points      79
```

```
few points          27
```

```
many points         25
```

```
very few points     12
```

```
dtype: int64
```

```
>> wnba['PTS_ordinal_scale'].value_counts().sort_index()
```

```
a lot of points      79
```

```
few points          27
```

```
many points         25
```

```
very few points     12
```

```
dtype: int64
```

# Sorting Tables for Ordinal Variables

---

```
>> wnba['PTS_ordinal_scale'].value_counts()
```

```
a lot of points      79
```

```
few points           27
```

```
many points          25
```

```
very few points      12
```

```
dtype: int64
```

```
>> wnba['PTS_ordinal_scale'].value_counts().iloc[[3, 1, 2, 0]]
```

```
very few points      12
```

```
few points           27
```

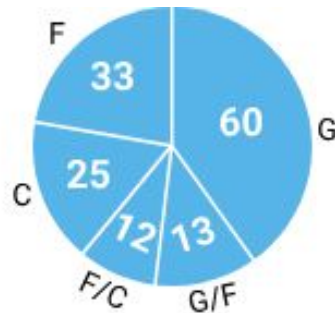
```
many points          25
```

```
a lot of points      79
```

```
dtype: int64
```

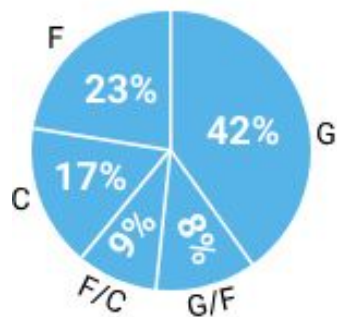
143 players

Counting players  
by position

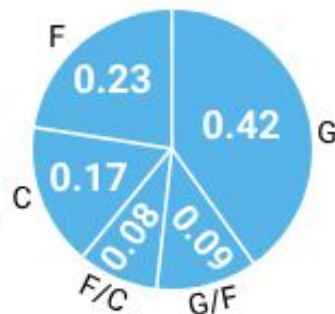


Divide frequencies by  
143 to get **proportions**\*

```
>> wnba['Pos'].value_counts(normalize = True) * 100  
G      41.958042  
F      23.076923  
C      17.482517  
G/F     9.090909  
F/C     8.391608
```



Multiply proportions  
by 100 to get **percentages**

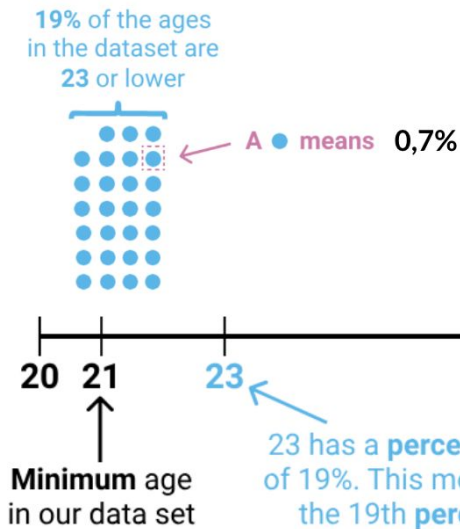


# Proportions and Percentages

# Percentiles and Percentile Ranks

```
1 percentages = wnba['Age'].value_counts(normalize = True).sort_index() * 100
2 percentages
```

```
21    1.398601
22    6.993007
23   10.489510
24   11.188811
25   10.489510
26    8.391608
27    9.090909
28    9.790210
29    5.594406
30    6.293706
31    5.594406
32    5.594406
33    2.097902
34    3.496503
35    2.797203
36    0.699301
```



"What percentage of players are 23 years or younger?"

# Percentiles and Percentile Ranks

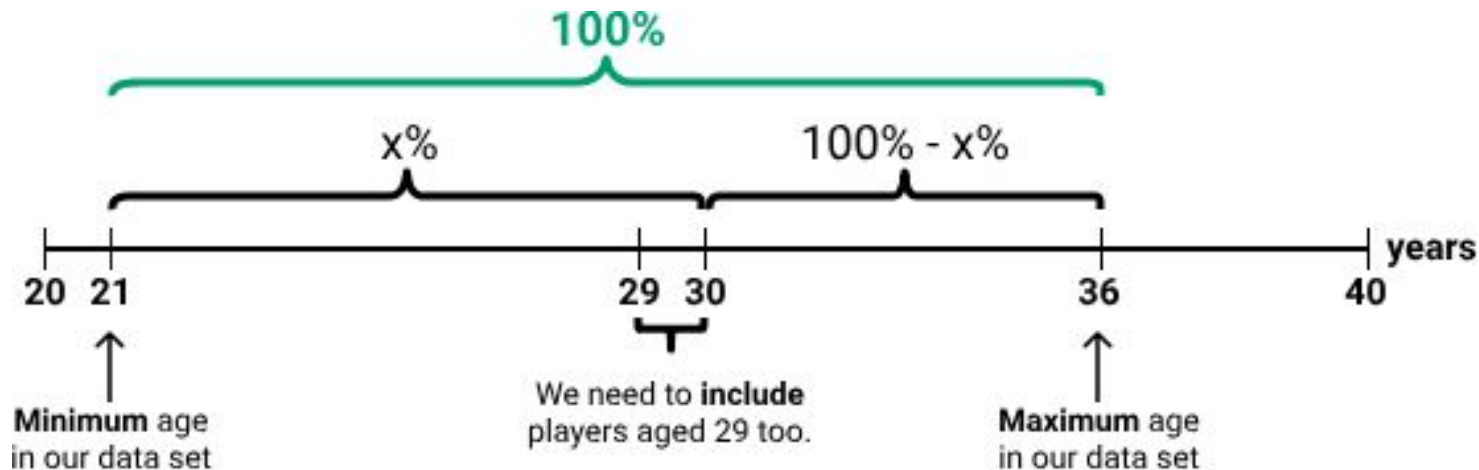
---

"What percentage of players are 23 years or younger?"

```
>> from scipy.stats import percentileofscore  
>> percentileofscore(a = wnba['Age'], score = 23, kind = 'weak')  
18.88111888111888
```

# Percentiles and Percentile Ranks

"What percentage of players are 30 years or older?"



```
>> 100 - percentileofscore(wnba['Age'], 29, kind = 'weak')  
26.573426573426573
```

# Finding Percentiles with Pandas

---

```
>> wnba['Age'].describe()
```

```
count    143.000000
```

```
mean      27.076923
```

```
std        3.679170
```

```
min        21.000000
```

```
25%        24.000000
```

```
50%        27.000000
```

```
75%        30.000000
```

```
max        36.000000
```

```
Name: Age, dtype: float64
```

```
>>wnba['Age'].describe().iloc[3:]
```

```
min        21.0
```

```
25%        24.0
```

```
50%        27.0
```

```
75%        30.0
```

```
max        36.0
```

```
Name: Age, dtype: float64
```

# Finding Percentiles with Pandas

```
>> wnba['Age'].describe(percentiles = [.1, .15, .33, .5, .592, .85, .9]).iloc[3:]
```

min	21.0
10%	23.0
15%	23.0
33%	25.0
50%	27.0
59.2%	28.0
85%	31.0
90%	32.0
max	36.0

Name: Age, dtype: float64

Percentile	Age
0%	21
25%	24
50%	27
75%	30
100%	36

25% of ages are 24 or less  
50% of ages are 27 or less  
75% of ages are 30 or less  
100% of ages are 36 or less

Minimum age in our data set

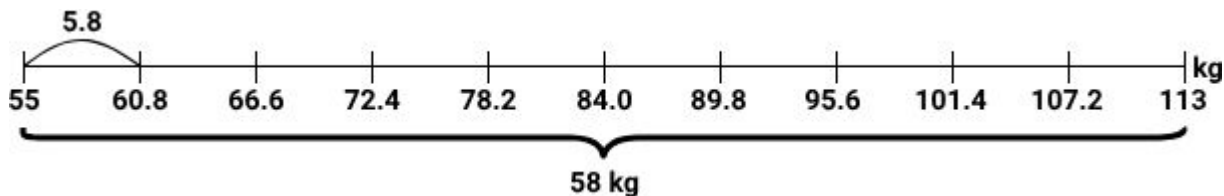
Maximum age in our data set



# Grouped Frequency Distribution Tables

```
>> wnba['Weight'].value_counts(bins = 10).sort_index()
```

(54.941, 60.8]	5
(60.8, 66.6]	21
(66.6, 72.4]	10
(72.4, 78.2]	33
(78.2, 84.0]	31
(84.0, 89.8]	24
(89.8, 95.6]	10
(95.6, 101.4]	3
(101.4, 107.2]	2
(107.2, 113.0]	3



# Information Loss

---

```
>> wnba['PTS'].value_counts(bins = 10)
```

```
(1.417, 60.2]      30
```

```
(60.2, 118.4]     24
```

```
(118.4, 176.6]    17
```

```
(176.6, 234.8]    20
```

```
(234.8, 293.0]    17
```

```
(293.0, 351.2]     8
```

```
(351.2, 409.4]    10
```

```
(409.4, 467.6]     8
```

```
(467.6, 525.8]     4
```

```
(525.8, 584.0]     5
```

```
Name: PTS, dtype: int64
```

```
wnba['PTS'].value_counts(bins = 5).sort_index()
```

```
(1.417, 118.4]      54
```


```
(118.4, 234.8]      37
```

```
(234.8, 351.2]      25
```

```
(351.2, 467.6]      18
```

```
(467.6, 584.0]       9
```

# Information Loss




# Readability for Grouped Frequency Tables

---

```
wnba[ 'PTS' ].value_counts(bins = 5).sort_index()
```

(1.417, 118.4]	54	(0, 100]	49
(118.4, 234.8]	37	(100, 200]	28
(234.8, 351.2]	25	(200, 300]	32
(351.2, 467.6]	18	(300, 400]	17
(467.6, 584.0]	9	(400, 500]	10
		(500, 600]	7

Name: PTS, dtype: int64



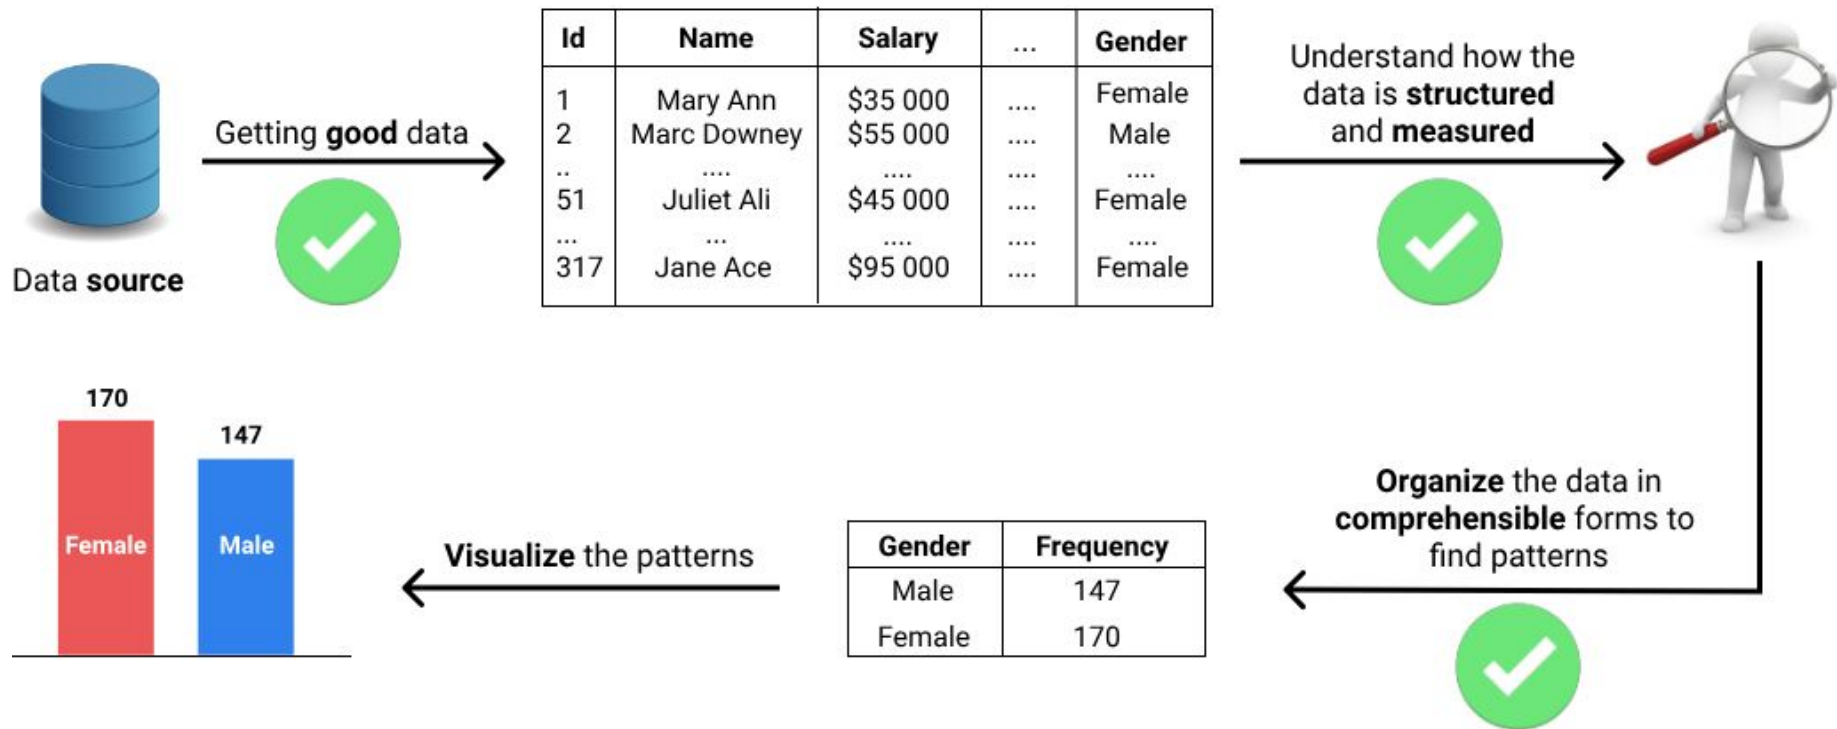
# Readability for Grouped Frequency Tables

```
>> intervals = pd.interval_range(start = 0, end = 600, freq = 100)
>> intervals
IntervalIndex([(0, 100], (100, 200], (200, 300], (300, 400], (400, 500], (500, 600]]
              closed='right',
              dtype='interval[int64]')

>> gr_freq_table = pd.Series([0,0,0,0,0,0], index = intervals)
>> gr_freq_table
(0, 100]      0
(100, 200]    0
(200, 300]    0
(300, 400]    0
(400, 500]    0
(500, 600]    0
dtype: int64

>> for value in wnba['PTS']:
    for interval in intervals:
        if value in interval:
            gr_freq_table.loc[interval] += 1
        break
```

```
>> gr_freq_table
(0, 100]      49
(100, 200]    28
(200, 300]    32
(300, 400]    17
(400, 500]    10
(500, 600]     7
```

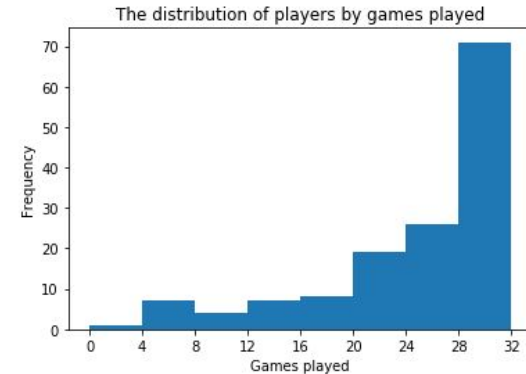

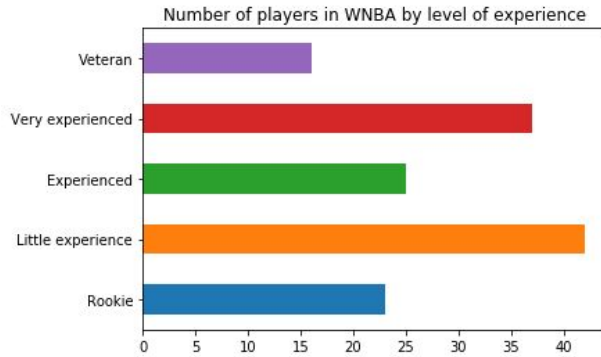


# Lesson 15 Statistical Fundamentals II.ipynb

## Section 1



# Visualizing Distributions




Graphs make easy to scan and compare frequencies, providing us with a single picture of the entire distribution of a variable (**nominal** or **ordinal scale**)




# Bar Plots

horizontal bar plots are ideal to use when the labels of the unique values are long




```
wnba['Pos'].value_counts().plot.bar()
```



```
wnba['Pos'].value_counts().plot.barh()
```


# Pie Charts



# Pie Charts


---

Percentage of players in WNBA by level of experience



```
wnba['Exp_ordinal'].value_counts().\nplot.pie(figsize = (6,6),\n          autopct = '%.2f%%',\n          title = 'Percentage of players in \\\nWNBA by level of experience')\nplt.ylabel('')
```


# Histograms



We can see that 75% of the values are distributed within a relatively narrow interval (between 2 and 277), while the remaining 25% are distributed in an interval that's slightly larger.

```
>> wnba['PTS'].describe()  
count    143.000000  
mean      201.790210  
std       153.381548  
min        2.000000  
25%       75.000000  
50%      177.000000  
75%      277.500000  
max      584.000000
```

# The Statistics Behind Histograms




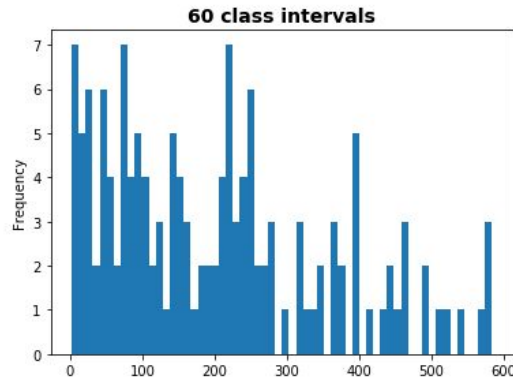

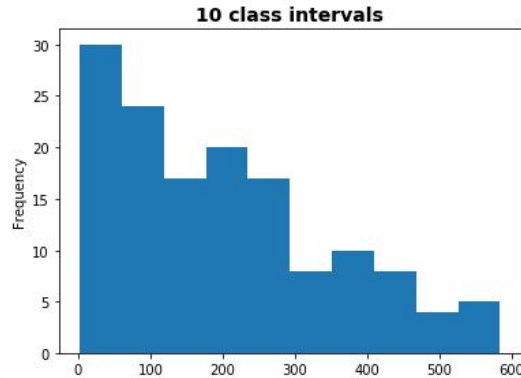

We can see immediately that roughly **three quarters (75%)** of the values are within this interval

The **remaining quarter** is within this interval


```
>> wnba['PTS'].describe()
count    143.000000
mean     201.790210
std      153.381548
min       2.000000
25%      75.000000
50%     177.000000
75%     277.500000
max     584.000000
Name: PTS, dtype: float64
```

```
>> wnba['PTS'].plot.hist()
```

# Binning for Histograms



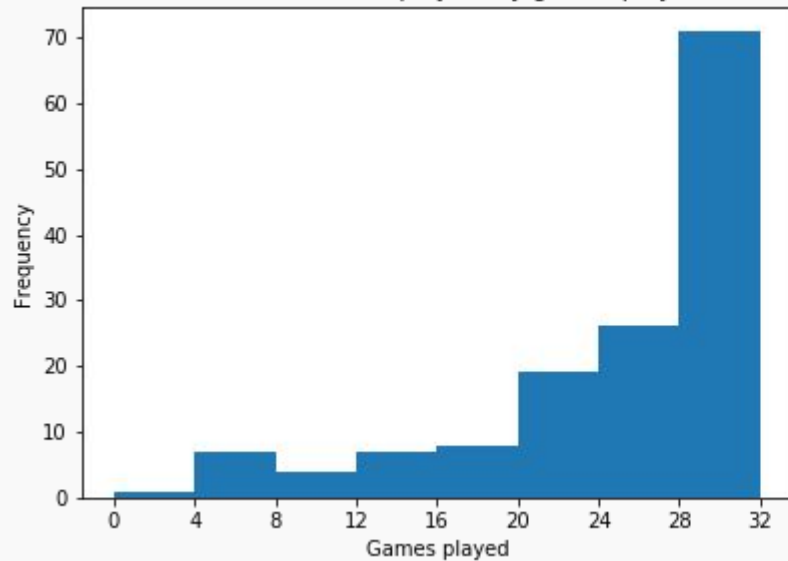
# Skewed Distributions




# Skewed Distributions

---

The distribution of players by games played




The distribution of players by total points


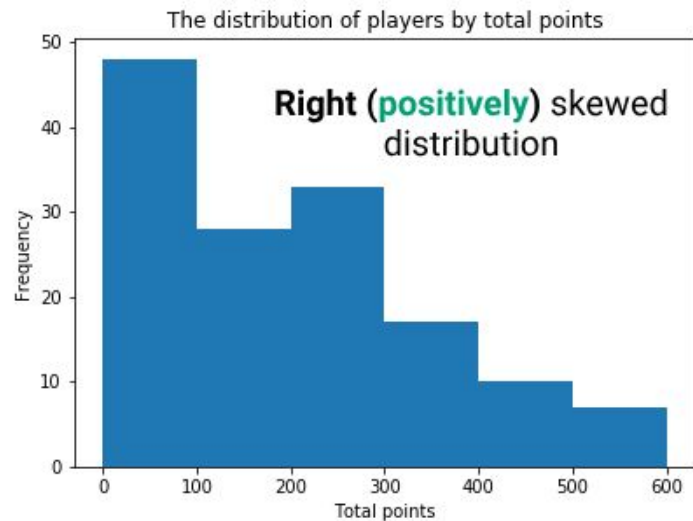




# Skewed Distributions

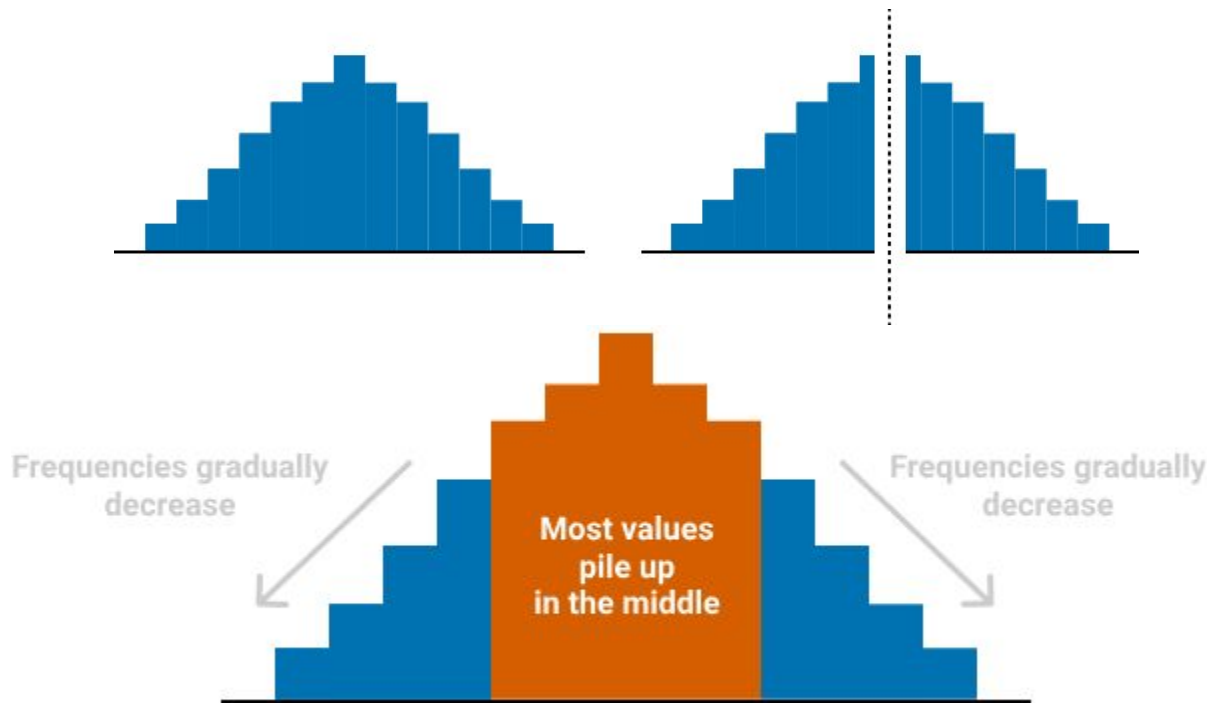


←  
If the **tail** points in the direction of **negative numbers** then the distribution is **negatively skewed**



→  
If the **tail** points in the direction of **positive numbers** then the distribution is **positively skewed**


# Symmetrical Distributions







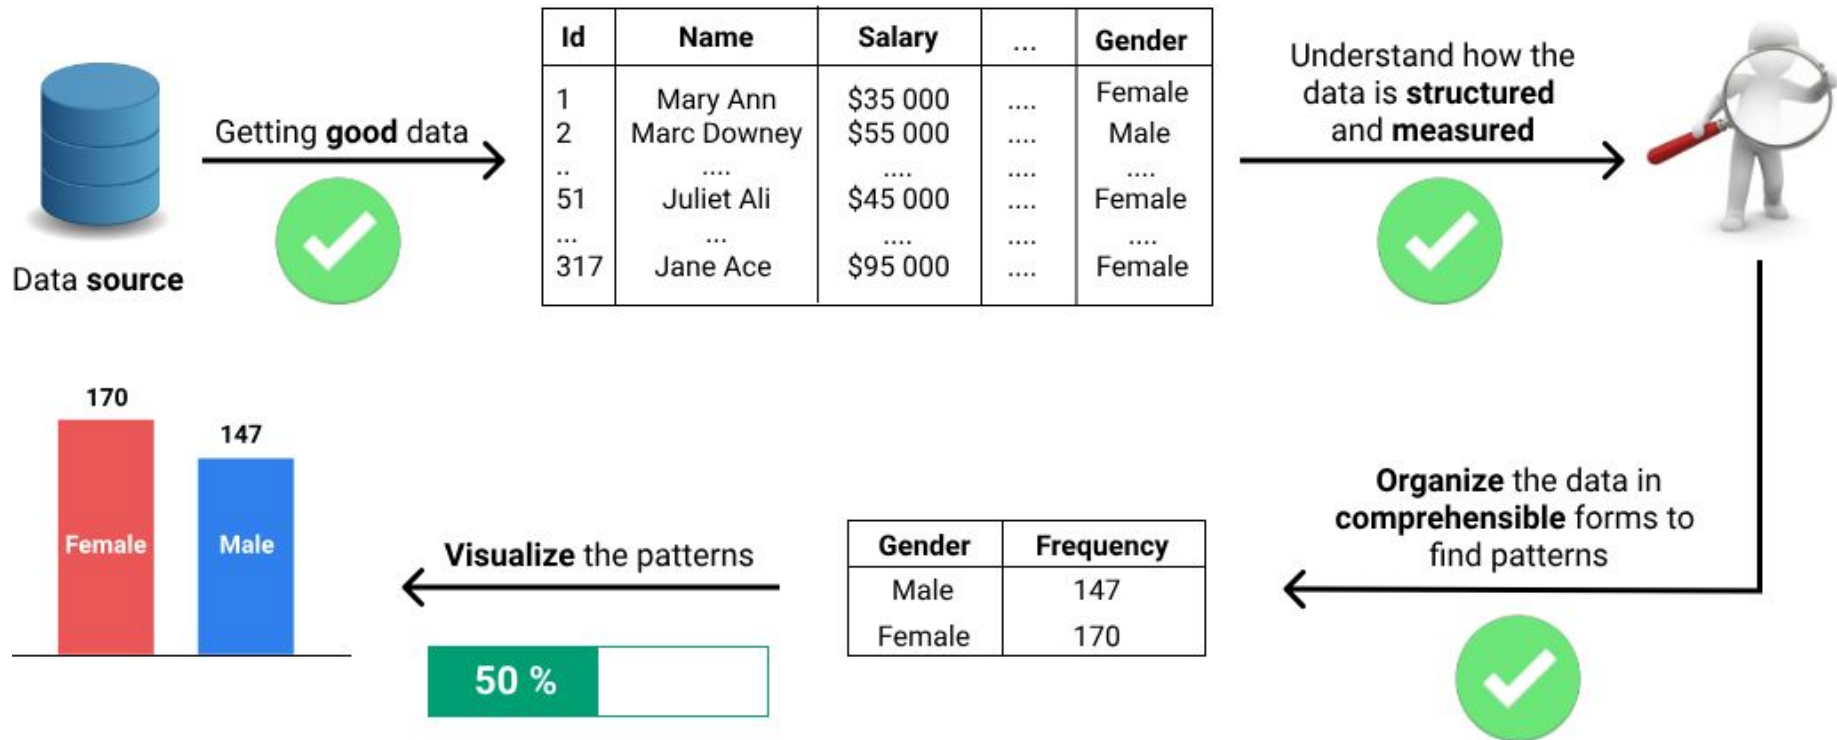
# Symmetrical Distribution (uniform)

---

The values are distributed uniformly



Scale of measurement	Graphs we can use to show the distribution
Nominal	 A bar chart with four bars of different colors (yellow, green, orange, blue) and a pie chart divided into five segments, representing nominal data.
Ordinal	 A bar chart with four bars of different colors (yellow, green, orange, blue) and a pie chart divided into five segments, representing ordinal data.
Interval	 A histogram with many blue bars forming a bell-shaped curve, representing interval data.
Ratio	 A histogram with many blue bars forming a bell-shaped curve, representing ratio data.




# Lesson 15 Statistical Fundamentals II.ipynb

## Section 2

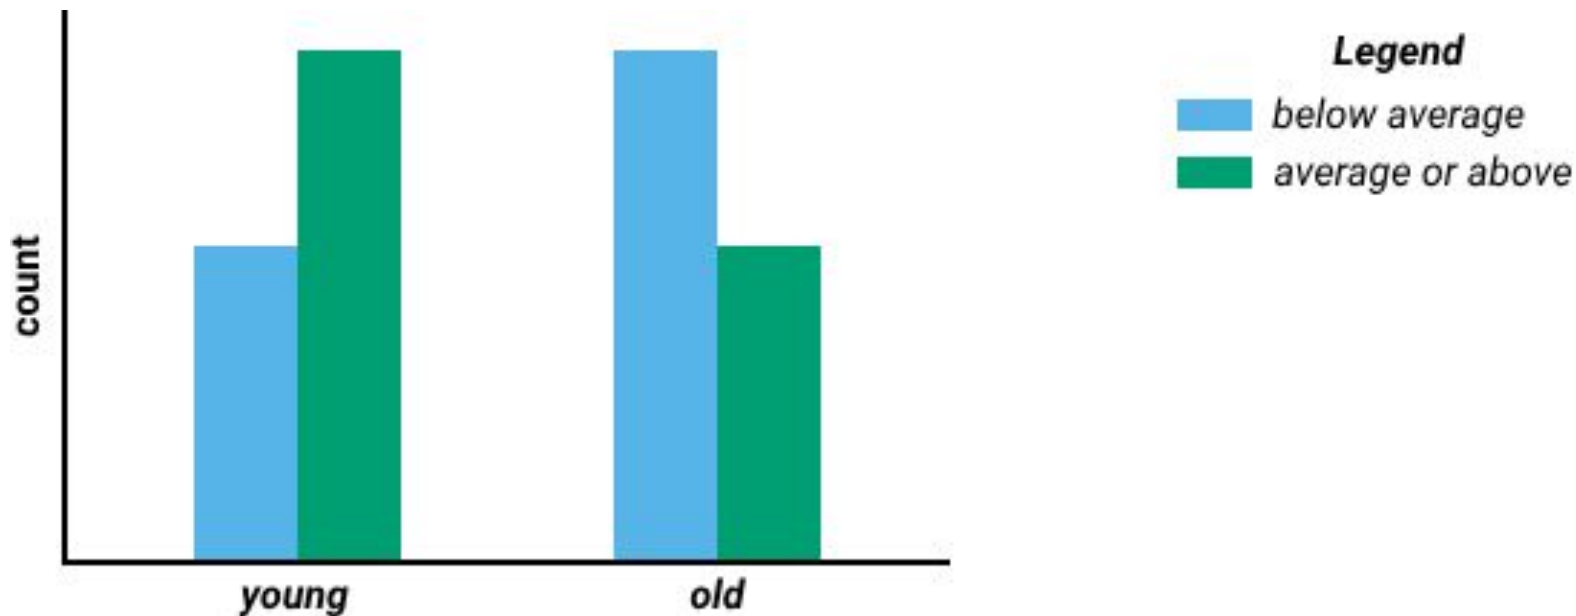


# Comparing Frequency Distribution




Years in WNBA	Label
0	Rookie
1-3	Little experience
4-5	Experienced
5-10	Very experienced
>10	Veteran

# Challenge: Do older players play less?









```
sns.countplot(x = 'age_mean_relative', hue = 'min_mean_relative', data = wnba)
```

# Comparing Histograms




```
wnba[wnba.Age >= 27]['MIN'].plot.hist(label = 'Old', legend = True)
wnba[wnba.Age < 27]['MIN'].plot.hist(label = 'Young', legend = True)
```




```
sns.set_style("white")
wnba[wnba.Age >= 27]['MIN'].plot.hist(histtype = 'step',
label = 'Old',
legend = True,color="red")
```

# Comparing Histograms




```
plt.axvline(497, label = 'Average')
```



# Kernel Density Estimate (KDE) Plots





# Kernel Density Estimate Plots





# Drawbacks of Kernel Density Plots



# Strip Plots




# Box Plots

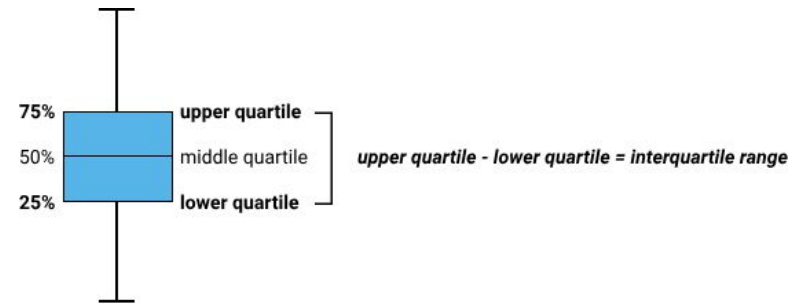






# Box Plots



# Outliers




# Outliers



```
>> wnba[wnba['Pos'] == 'C']['Height'].describe()
count      25.000000
mean       194.920000
std         4.132392
min        185.000000
25%        193.000000
50%        193.000000
75%        196.000000
max        206.000000
Name: Height, dtype: float64
```


# Outliers



# Outliers


```
sns.boxplot(wnba[wnba['Pos'] == 'C']['Height'], whis = 1.5,  
            orient = 'vertical', width = .45)
```



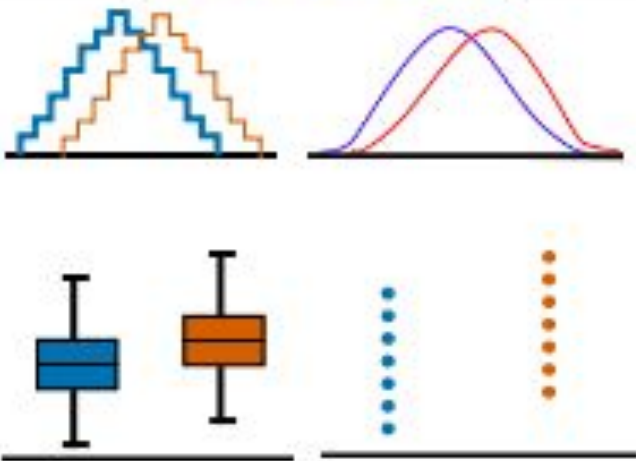
<matplotlib.axes.\_subplots.AxesSubplot at 0x1a180c4518>



```
sns.boxplot(wnba[wnba['Pos'] == 'C']['Height'], whis = 4,  
            orient = 'vertical', width = .45)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x1a18180208>



Scale of measurement	Graphs we can use to compare distributions
Nominal	 A bar chart with six categories on the x-axis. Each category has three bars of different colors (yellow, pink, blue). The heights of the bars vary across categories, representing the frequency of each category.
Ordinal	 A bar chart with six categories on the x-axis. Each category has three bars of different colors (yellow, pink, blue). The heights of the bars vary across categories, representing the frequency of each category.
Interval & Ratio	 Two sets of graphs. The top set shows two overlapping bell curves, one blue and one orange, representing normal distributions. The bottom set shows two box plots, one blue and one orange, representing the distribution of data. The blue box plot is on the left and the orange box plot is on the right. The y-axis for the box plots is marked with dots.

