

Data Science & ML Course

Lesson #17 K-Nearest Neighbors

Ivanovitch Silva
November, 2018



Update from repository

```
git clone https://github.com/ivanovitchm/datascience2machinelearning.git
```

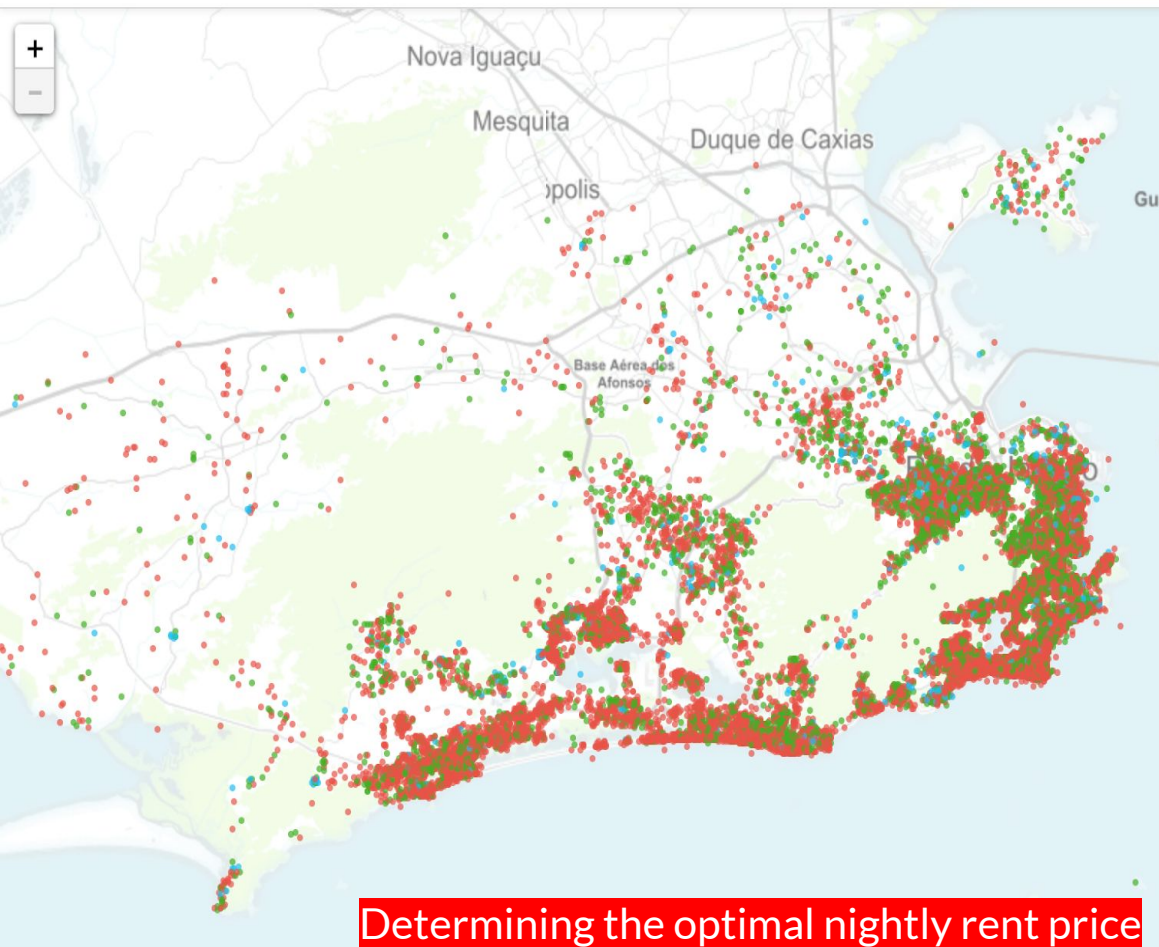
Or

```
git pull
```



Agenda

- Univariate KNN
 - Euclidean distance for univariate
 - Function to make predictions
 - Error metrics
- Multivariate KNN
 - Normalize columns
 - Euclidean distance for multivariate
- Hyperparameter optimization
- Cross-Validation



Determining the optimal nightly rent price

Rio de Janeiro

Filter by:

Rio de Janeiro

32,469

out of 32,469 listings (100%)

[About Airbnb in Rio de Janeiro](#)

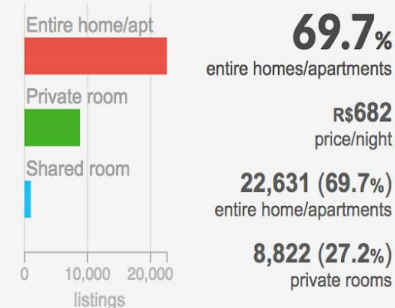
How is Airbnb really being used in and affecting your neighbourhoods?

Room Type

☐ Only entire homes/apartments

Airbnb hosts can list entire homes/apartments, private or shared rooms.

Depending on the room type and [activity](#), an airbnb listing could be more like a hotel, disruptive for neighbours, taking away housing, and [illegal](#).



69.7%
entire homes/apartments

R\$682
price/night

22,631 (69.7%)
entire home/apartments

8,822 (27.2%)
private rooms


1,016 (3.1%)
shared rooms

- **host_response_rate**: the response rate of the host
- **host_acceptance_rate**: number of requests to the host that convert to rentals
- **host_listings_count**: number of other listings the host has
- **latitude**: latitude dimension of the geographic coordinates
- **longitude**: longitude part of the coordinates
- **city**: the city the living space resides
- **zipcode**: the zip code the living space resides
- **state**: the state the living space resides
- **accommodates**: the number of guests the rental can accommodate
- **room_type**: the type of living space (Private room, Shared room or Entire home/apt)
- **bedrooms**: number of bedrooms included in the rental
- **bathrooms**: number of bathrooms included in the rental
- **beds**: number of beds included in the rental
- **price**: nightly price for the rental
- **cleaning_fee**: additional fee used for cleaning the living space after the guest leaves
- **security_deposit**: refundable security deposit, in case of damages
- **minimum_nights**: minimum number of nights a guest can stay for the rental
- **maximum_nightss**: maximum number of nights a guest can stay for the rental
- **number_of_reviews**: number of reviews that previous guests have left

96
cols

```
total 112M
-rw-r--r-- 1 root root 112M Sep 29 21:46 listings.csv
drwxr-xr-x 2 root root  1M Sep 27 20:32 sample_data
```


Select the number of similar listings, **k**, you want to compare with.



For each listing, calculate how similar it is to our unpriced listing.



For this example, we'll use 3 for our **k** value.

K-Nearest Neighbors

Rank each listing by the similarity metric and select the first **k** listings.

dataset (ordered by similarity)		
bedrooms	price	similarity
1	160	0
1	60	0
1	95	0
1	50	0
3	350	2

our unpriced listing	
bedrooms	price
1	?

Calculate the mean list price for the **k** similar listings and use as our list price.



Euclidean distance

$$d = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2}$$

index	host_listings_count	accommodates	bedrooms	bathrooms	beds
0	26	4	1	1	2
1	1	6	3	3	3

$$(q_1 - p_1) + (q_2 - p_2) + \dots + (q_n - p_n) \quad \text{differences} \quad (26 - 1) + (4 - 6) + (1 - 3) + (1 - 3) + (2 - 3)$$

$$(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2 \quad \text{squared differences} \quad (25)^2 + (-2)^2 + (-2)^2 + (-2)^2 + (-1)^2$$

$$\sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2}$$

Euclidean distance	= $\sqrt{625 + 4 + 4 + 4 + 1}$
	= $\sqrt{638}$
	= 25.258661

Euclidean distance - example

Univariate case

$$d = \sqrt{(q_1 - p_1)^2}$$


$$d = |q_1 - p_1|$$

	accommodates
our listing	8

	index	accommodates	distance
dc_listings	0	4	$(4 - 8)^2$
	1	6	$(6 - 8)^2$
	2	1	$(1 - 8)^2$
	3	2	$(2 - 8)^2$

Testing quality of predictions

Machine Learning Model



Error metrics

Mean Absolute Error

$$MAE = \frac{|actual_1 - predicted_1| + |actual_2 - predicted_2| + \dots + |actual_n - predicted_n|}{n}$$

Mean Squared Error

$$MSE = \frac{(actual_1 - predicted_1)^2 + (actual_2 - predicted_2)^2 + \dots + (actual_n - predicted_n)^2}{n}$$

Root Mean Squared Error

$$RMSE = \sqrt{MSE}$$

Improve the accuracy (multivariate knn)

- **Increase the number of attributes** the model uses to calculate similarity when ranking the closest neighbors
- **Increase k**, the number of nearby neighbors the model uses when computing the prediction
- We'll focus on more robust techniques for testing a machine learning model's accuracy, namely, **cross-validation**

Improve the accuracy (multivariate knn)

When selecting more attributes to use in the model, we need to watch out for columns that don't work well with the distance equation.

- **non-numerical values** (e.g. city or state)
 - Euclidean distance equation expects numerical values
- **missing values**
 - distance equation expects a value for each observation and attribute
- **non-ordinal values** (e.g. latitude or longitude)
 - ranking by Euclidean distance doesn't make sense if all attributes aren't ordinal

Removing features

- `room_type`: e.g. **Private room**
- `city`: e.g. **Washington**
- `state`: e.g. **DC**
- `host_response_rate`
- `host_acceptance_rate`
- `host_listings_count`
- `latitude`: e.g. **38.913458**
- `longitude`: e.g. **-77.031**
- `zipcode`: e.g. **20009**

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 37245 entries, 3629 to 33003
Data columns (total 10 columns):
accommodates      37245 non-null int64
bathrooms         37168 non-null float64
bedrooms          37219 non-null float64
beds              37188 non-null float64
price             37245 non-null float64
security_deposit   18540 non-null object
cleaning_fee      23159 non-null object
minimum_nights    37245 non-null int64
maximum_nights    37245 non-null int64
number_of_reviews  37245 non-null int64
dtypes: float64(4), int64(4), object(2)
memory usage: 3.1+ MB
```

Normalize columns

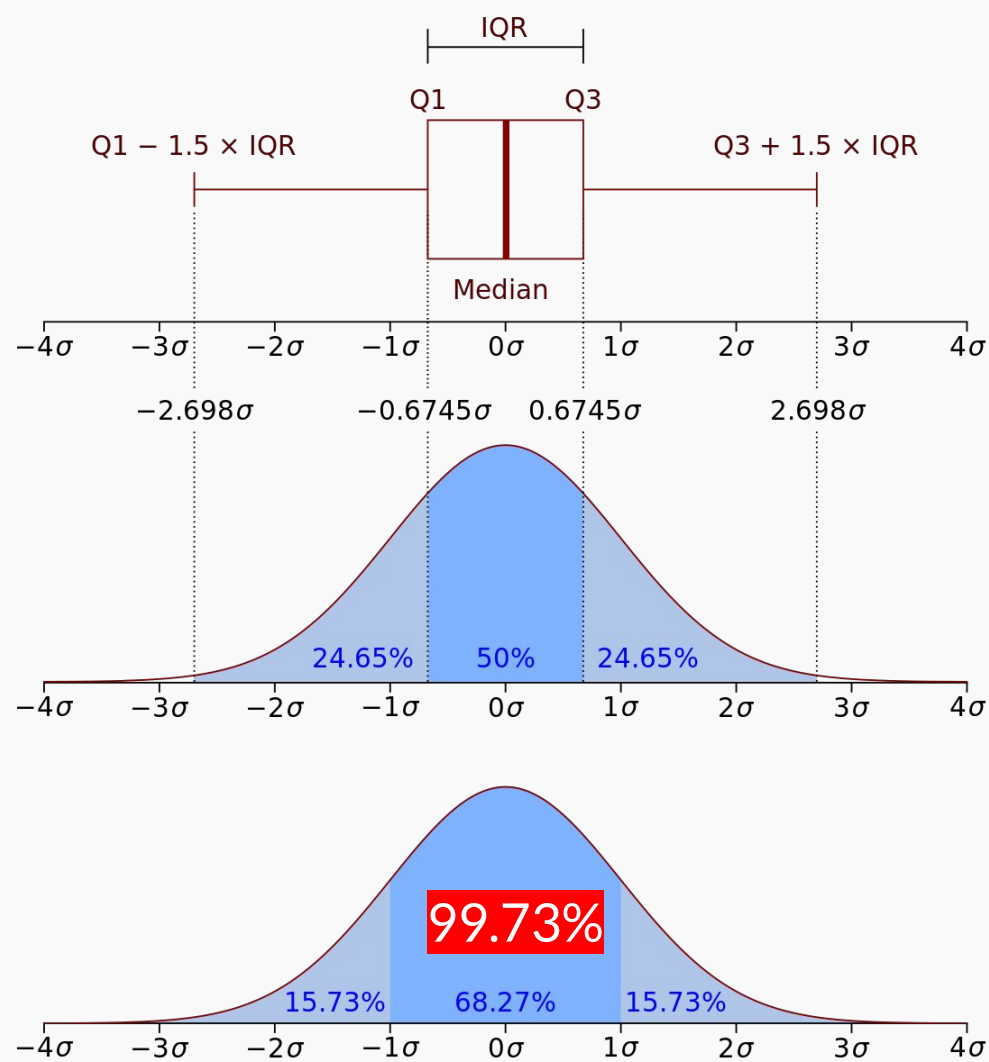
accommodates	bathrooms	bedrooms	beds	price	minimum_nights	maximum_nights	number_of_reviews
8	1.0	1.0	1.0	200.0	2	60	0
6	2.0	3.0	4.0	901.0	10	27	0
2	1.0	1.0	1.0	229.0	2	50	24
2	1.0	1.0	2.0	200.0	3	1125	0
2	2.0	1.0	1.0	75.0	2	30	2

accommodates	bathrooms	bedrooms	beds	price	minimum_nights	maximum_nights	number_of_reviews
1.455750	-0.67557	-0.597149	-0.779645	200.0	-0.121917	-0.008480	-0.337027
0.691698	0.27004	1.229023	0.665006	901.0	0.254259	-0.008483	-0.337027
-0.836405	-0.67557	-0.597149	-0.779645	229.0	-0.121917	-0.008481	0.981934
-0.836405	-0.67557	-0.597149	-0.298095	200.0	-0.074895	-0.008400	-0.337027
-0.836405	0.27004	-0.597149	-0.779645	75.0	-0.121917	-0.008482	-0.227114

Normalize columns (outliers)

```
1 normalized_listings.describe().
```

	accommodates	bathrooms	bedrooms	beds	price	minimum_nights	maximum_nights	number_of_reviews
count	3.712900e+04	3.712900e+04	3.712900e+04	3.712900e+04	37129.000000	3.712900e+04	3.712900e+04	3.712900e+04
mean	-1.314721e-16	-4.592913e-18	-4.822558e-17	-4.478090e-17	645.284872	-2.066811e-17	7.654855e-19	2.387358e-17
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1658.580291	1.000000e+00	1.000000e+00	1.000000e+00
min	-1.218415e+00	-1.621158e+00	-1.510215e+00	-1.261178e+00	0.000000	-1.689371e-01	-8.484452e-03	-3.370226e-01
25%	-8.363938e-01	-6.755609e-01	-5.971412e-01	-7.796344e-01	150.000000	-1.689371e-01	-8.482279e-03	-3.370226e-01
50%	-7.235234e-02	-6.755609e-01	-5.971412e-01	-2.980906e-01	296.000000	-1.219157e-01	-8.400222e-03	-3.370226e-01
75%	3.096684e-01	2.700359e-01	3.159325e-01	1.834533e-01	626.000000	-2.787293e-02	-8.400222e-03	-1.721546e-01
max	5.952288e+01	1.729078e+01	3.957810e+01	2.281602e+01	41719.000000	5.733824e+01	1.609190e+02	1.647951e+01



Euclidean distance for multivariate case

accommodates

bathrooms

-0.596544

-0.439151

-0.596544

0.412923

$$(q_1 - p_1) + (q_2 - p_2) + \dots + (q_n - p_n)$$

differences

$$(-0.596544 + 0.596544) + (-0.439151 - 0.412923)$$

$$(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2$$

squared
differences

$$(0)^2 + (-0.852074)^2$$

$$\sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2}$$

Euclidean
distance

$$= \sqrt{0 + 0.72603}$$

$$= 0.852074$$

Introduction to scikit-learn



scikit-learn
Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ... — Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ... — Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ... — Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization. — Examples

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics. — Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.

Modules: preprocessing, feature extraction. — Examples

Scikit-learn workflow

The scikit-learn workflow consists of 4 main steps:

- instantiate the specific machine learning model you want to use
- fit the model to the training data
- use the model to make predictions
- evaluate the accuracy of the predictions

Scikit-learn workflow

```
# features
train_columns = ['accommodates', 'bathrooms']

# instantiate a knn object
knn = KNeighborsRegressor(n_neighbors=5, algorithm='brute', metric='euclidean')

# train the model
knn.fit(train_df[train_columns], train_df['price'])

# predict|
predictions = knn.predict(test_df[train_columns])

# evaluate
rmse = np.sqrt(mean_squared_error(predictions, test_df.price))
```


Improve the accuracy (multivariate knn)

- We'll focus on the impact of **increasing k** , the number of nearby neighbors the model uses to make predictions.

Hyperparameters

- When we **vary the features** that are used in the model, we're **affecting the data** that the model uses.
- On the other hand, **varying the k value affects the behavior of the model independently of the actual data** that's used when making predictions.
- Values that affect the behavior and performance of a model that are unrelated to the data that's used are referred to as **hyperparameters**.

Hyperparameter Optimization

A simple but common [hyperparameter optimization](#) technique is known as [grid search](#):


- selecting a subset of the possible hyperparameter values,
- training a model using each of these hyperparameter values,
- evaluating each model's performance,
- selecting the hyperparameter value that resulted in the lowest error value.

Improve the accuracy (multivariate knn)

- We can focus on **more robust techniques for testing** a machine learning model's accuracy

Holdout Validation

Holdout Validation



Error

123.64


123.21

Mean Error


123.43

Holdout validation is actually a specific example of a larger class of validation techniques called **k-fold cross-validation**.

K-Fold Cross Validation



Bias-Variance Tradeoff





```
index.js
import React, { useState } from 'react';
import './index.css';

function App() {
  const [contacts, setContacts] = useState([
    { name: 'John Doe', phone: '123-456-7890' },
    { name: 'Jane Smith', phone: '987-654-3210' },
  ]);

  const handleClick = () => {
    // TODO: Implement handleClick logic
  };

  return (
    <div>
      <h1>Contact List</h1>
      <ul>
        {contacts.map(contact => (
          <li>{contact.name} {contact.phone}</li>
        ))}
      </ul>
      <button onClick={handleClick}>Add Contact</button>
    </div>
  );
}

export default App;
```

```
index.html
<!DOCTYPE html>
<html>
  <head>
    <script src="index.js"></script>
  </head>
  <body>
    <div>
      <h1>Contact List</h1>
      <ul>
        <li>John Doe 123-456-7890</li>
        <li>Jane Smith 987-654-3210</li>
      </ul>
      <button>Add Contact</button>
    </div>
  </body>
</html>
```