

Data Science & ML Course

Lesson #22 Random Forest

Ivanovitch Silva
December, 2018



Update from repository

```
git clone https://github.com/ivanovitchm/datascience2machinelearning.git
```

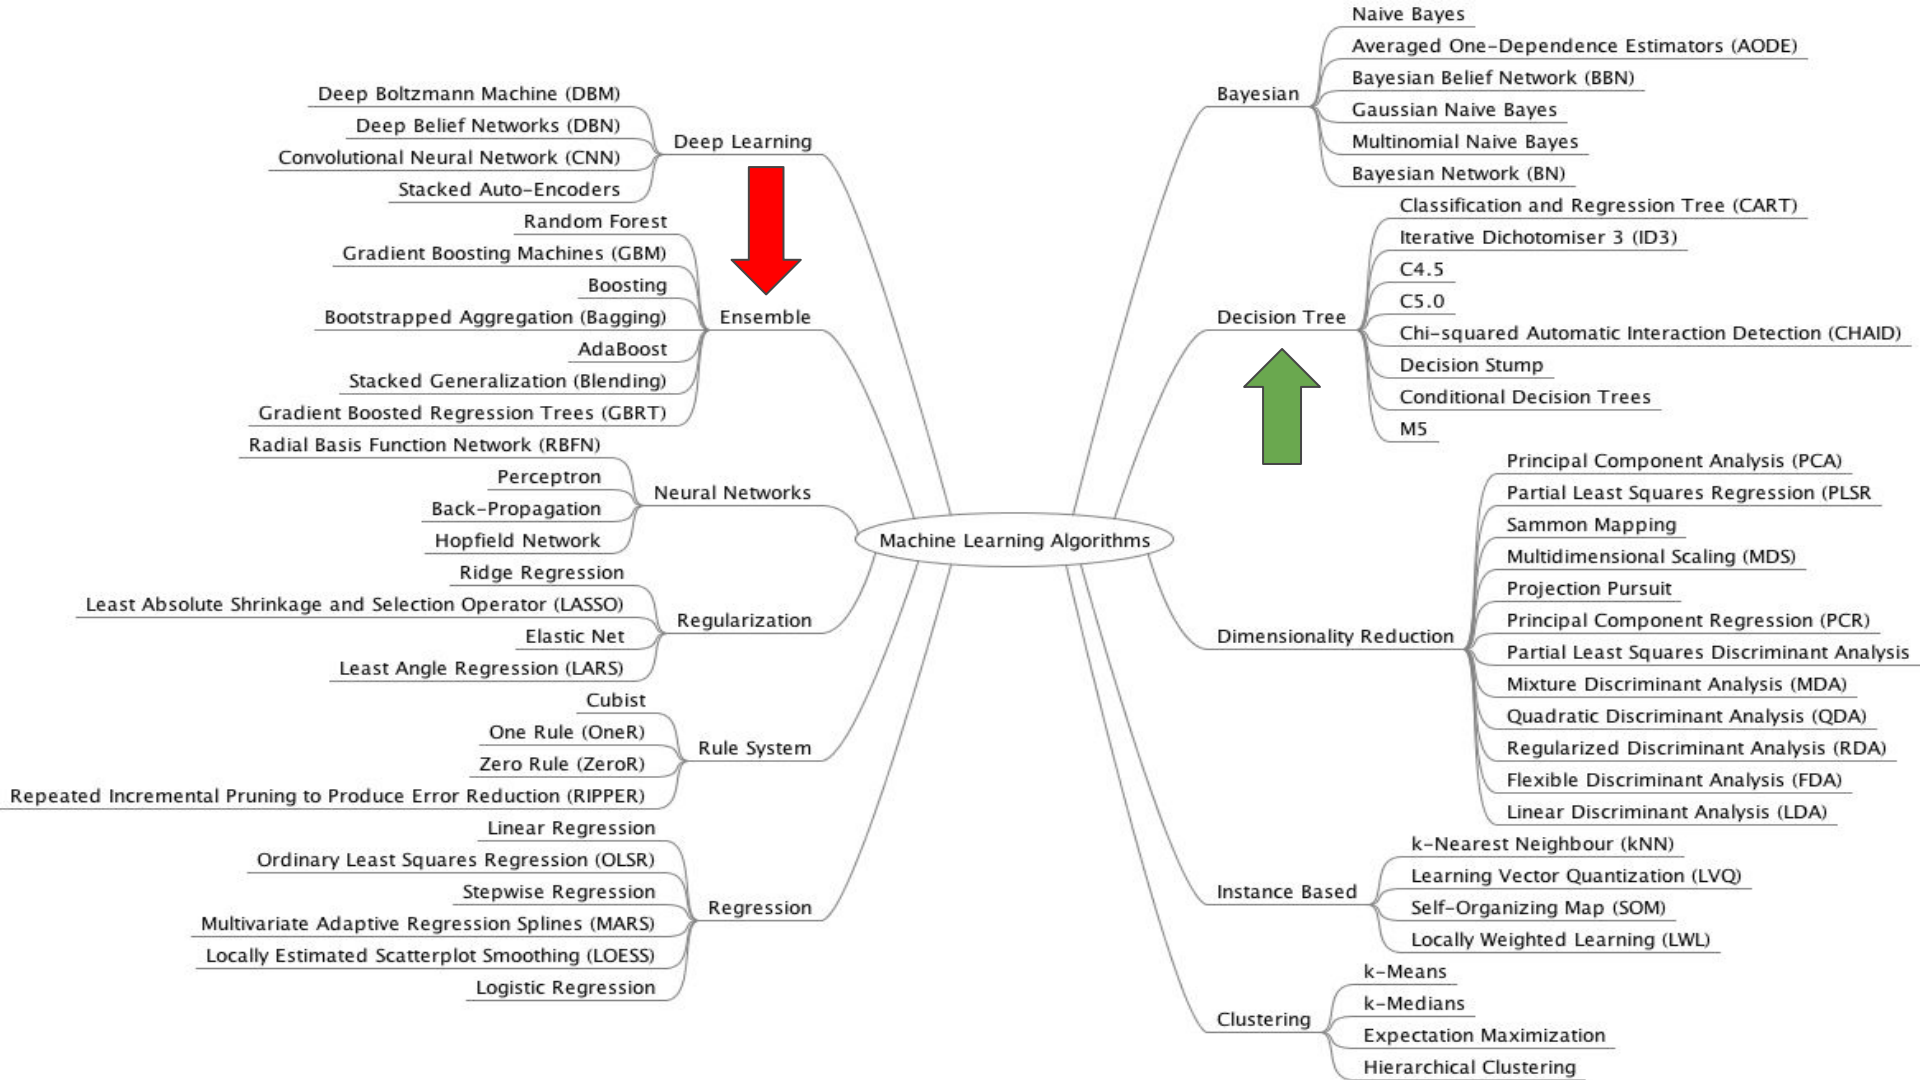
Or

```
git pull
```



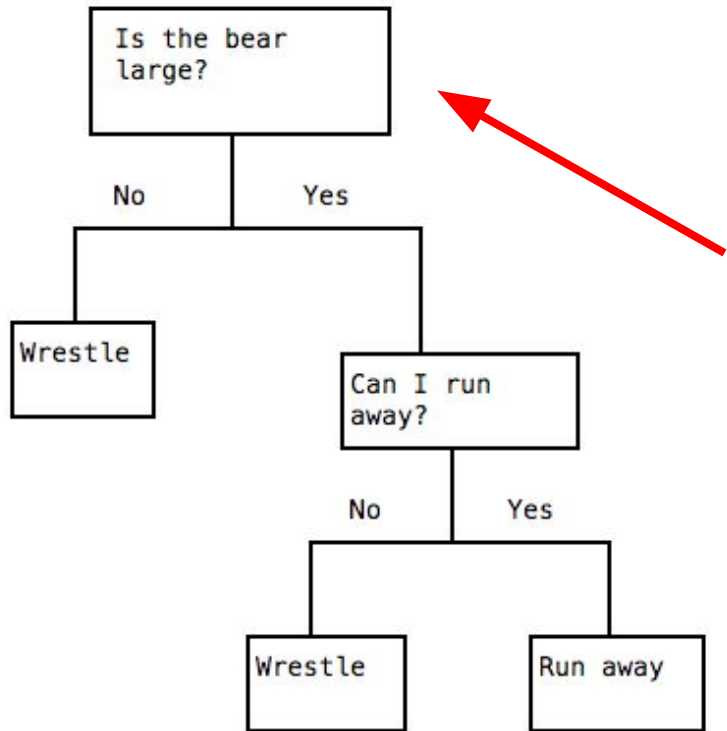
Agenda

- Previously on last class (Decision Trees)
- Ensembles (Random Forest)
- Combining predictions
- Why Ensembling works
- Introduction variation with bagging and random features
- Reducing overfitting using Random Forest
- Case study: US Census, predicting bike rentals




Decision Tree (classification)

Should I wrestle this bear?



Bear name	Size	Escape possible?	Action
Yogi	Small	No	Wrestle
Winnie	Small	Yes	Wrestle
Baloo	Large	Yes	Run away
Gentle Ben	Large	No	Wrestle

Receiver Operating Characteristic (ROC)

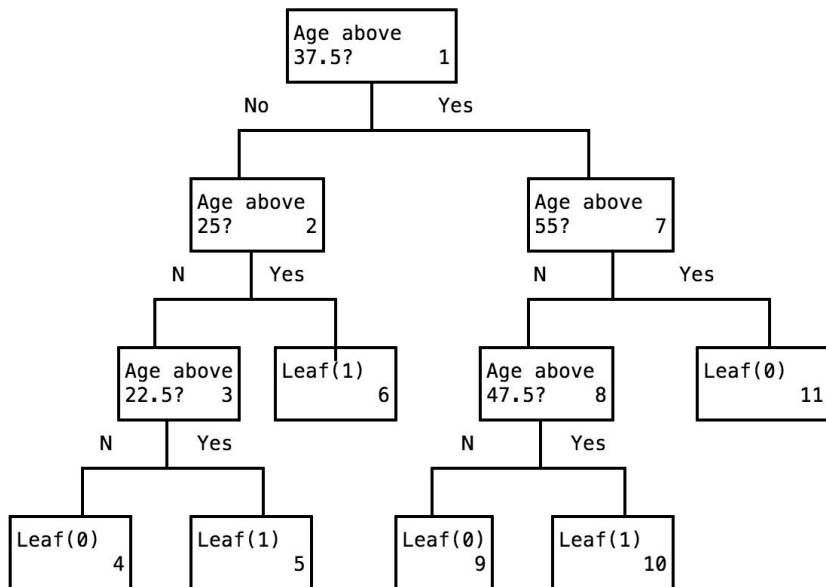


		True condition	
Total population		Condition positive	Condition negative
Predicted condition	Predicted condition positive	True positive, Power	False positive, Type I error
	Predicted condition negative	False negative, Type II error	True negative
		True positive rate (TPR), Recall, Sensitivity, probability of detection $= \frac{\sum \text{True positive}}{\sum \text{Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm $= \frac{\sum \text{False positive}}{\sum \text{Condition negative}}$

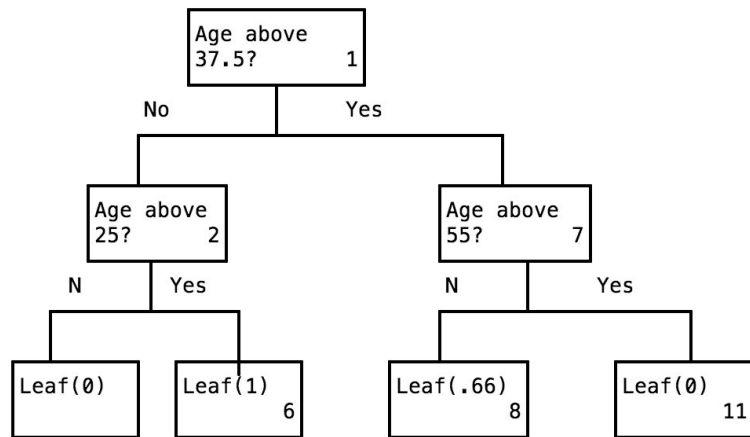
AUC - Area Under Curve

Decision Tree Overfitting


Full tree



Smaller tree



settings	train AUC	test AUC
default (min_samples_split: 2, max_depth: None)	0.947	0.694
min_samples_split: 13	0.842	0.699
min_samples_split: 13, max_depth: 7	0.748	0.743
min_samples_split: 100, max_depth: 2	0.662	0.655



Random Forest

Wisdom of the crowd



Logistic
Regression



SVM
Classifier




Random
Forest Classifier

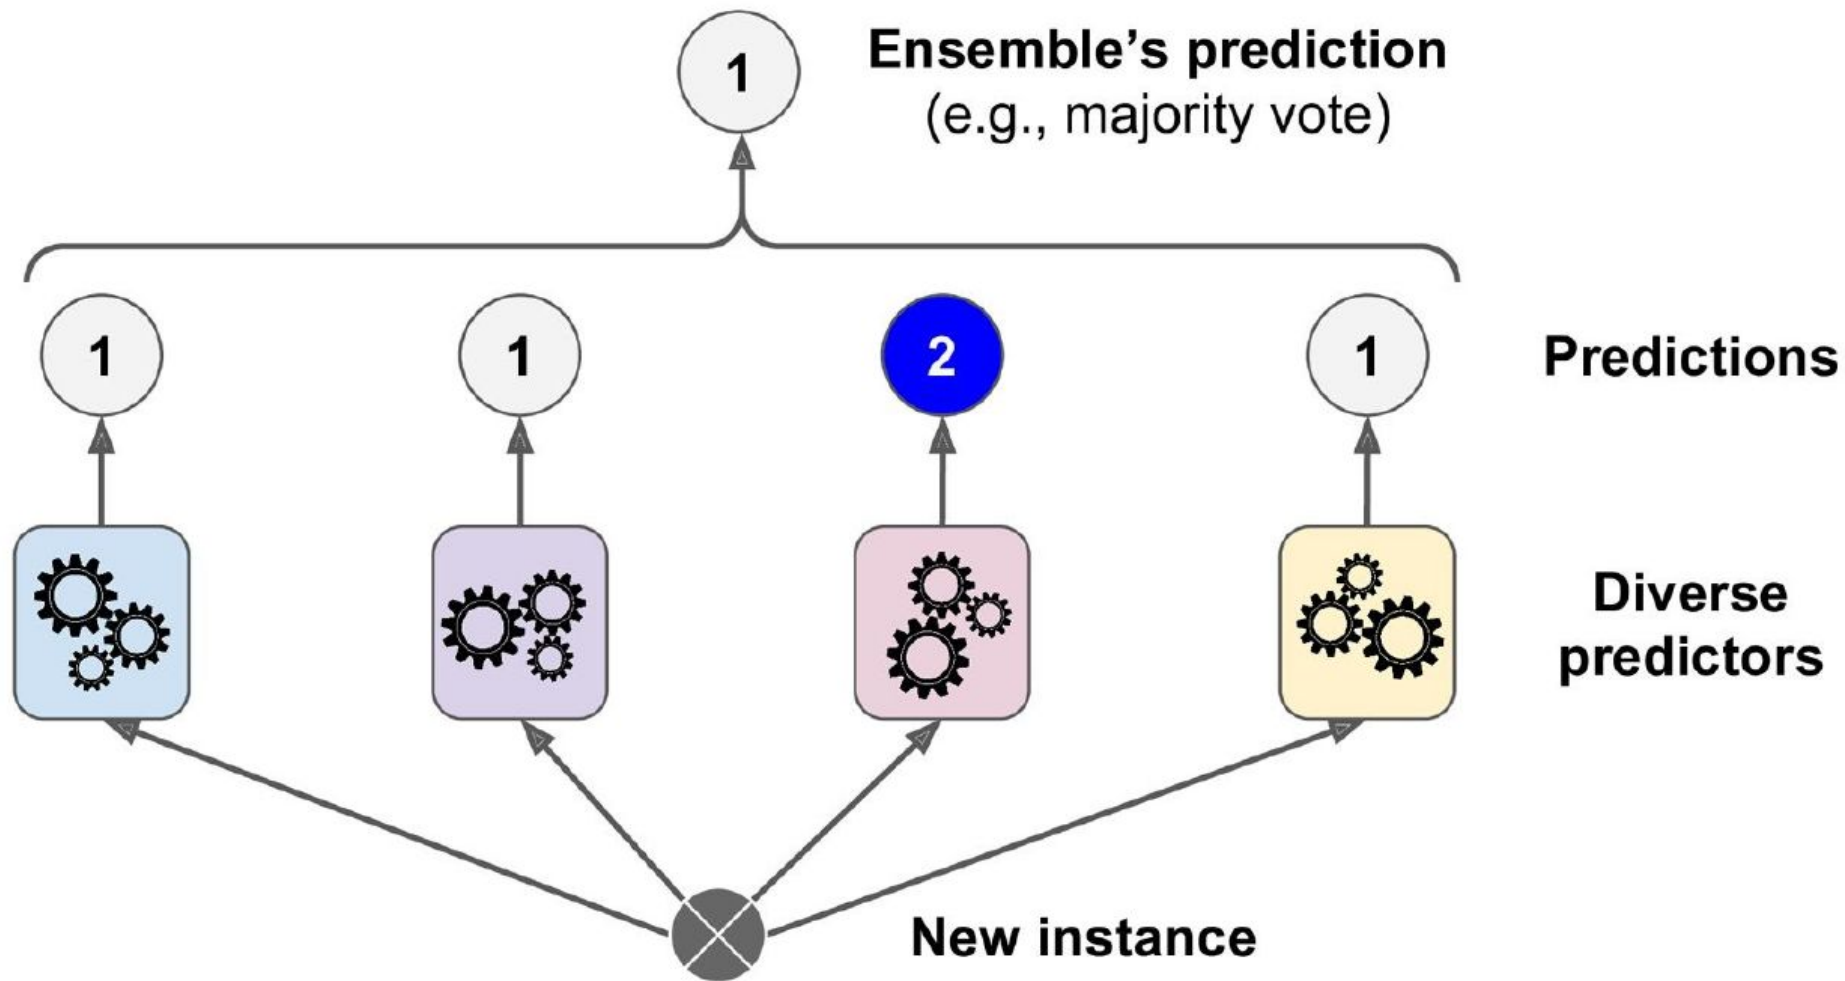


Other...



**Diverse
predictors**





Case Study #1: US Census

The target column, or what we want to predict, is whether individuals make less than or equal to 50k a year, or more than 50k a year.

US Census
1994

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0

Combining Model Predictions with Ensembles

```
1 from sklearn.tree import DecisionTreeClassifier
2 from sklearn.metrics import roc_auc_score
3
4 # features
5 columns = ["age", "workclass", "education_num", "marital_status",
6            "occupation", "relationship", "race", "sex",
7            "hours_per_week", "native_country"]
8
9 # model 1
10 clf = DecisionTreeClassifier(random_state=1, min_samples_leaf=2)
11 clf.fit(train[columns], train["high_income"])
12
13 # model 2
14 clf2 = DecisionTreeClassifier(random_state=1, max_depth=5)
15 clf2.fit(train[columns], train["high_income"])
16
17 # prediction on model 1
18 predictions = clf.predict(test[columns])
19 print(roc_auc_score(test["high_income"], predictions))
20
21 # prediction on model 2
22 predictions = clf2.predict(test[columns])
23 print(roc_auc_score(test["high_income"], predictions))
```

0.6878964226062301

0.6759853906508785

Combining our predictions

Majority Voting

DT1	DT2	DT3	Final Prediction
0	1	0	0
1	1	1	1
0	0	1	0
1	0	0	0

settings	test AUC
min_samples_leaf: 2	0.688
max_depth: 2	0.676
combined predictions	0.715

	#1 Model	#2 Model	Mean	Rounded
0	0.166667	0.288507	0.227587	0.0
1	0.000000	0.288507	0.144253	0.0
2	0.000000	0.180918	0.090459	0.0
3	0.000000	0.354167	0.177083	0.0
4	0.000000	0.041009	0.020504	0.0
5	0.000000	0.006875	0.003437	0.0
6	0.000000	0.006875	0.003437	0.0
7	0.333333	0.146179	0.239756	0.0
8	0.000000	0.006875	0.003437	0.0
9	0.666667	0.777431	0.722049	1.0
10	0.000000	0.041009	0.020504	0.0


Probability Voting

Why Ensembling Works

```
Terminal
  GOAT
  Successfully authenticated!
  ? Enter a name for the repository: example
  ? Optionally enter a description of the repository: Just an example
  ? Public or private: public
  ? Select the files and/or folders you wish to ignore: (Press <space> to select)
  > bower_components
    o index.js
    o lib
    o node_modules
    o package.json
```



Bagging and Pasting



Introduction Variation with Bagging

```

1 # We'll build 10 trees
2 tree_count = 10
3
4 # Each "bag" will have 60% of the number of original rows
5 bag_proportion = .6
6
7 predictions = []
8 for i in range(tree_count):
9     # We select 60% of the rows from train, sampling with replacement
10    # We set a random state to ensure we'll be able to replicate our results
11    # We set it to i instead of a fixed value so we don't
12    # get the same sample in every loop.
13    bag = train.sample(frac=bag_proportion, replace=True, random_state=i)
14
15    # Fit a decision tree model to the "bag"
16    clf = DecisionTreeClassifier(random_state=1, min_samples_leaf=2)
17    clf.fit(bag[colums], bag["high_income"])
18
19    # Using the model, make predictions on the test data
20    predictions.append(clf.predict_proba(test[colums])[:,1])
21 combined = np.sum(predictions, axis=0) / 10
22 rounded = np.round(combined)
23
24 print(roc_auc_score(test["high_income"], rounded))

```

0.7329963297474371

settings	test AUC
min_samples_leaf: 2	0.688
max_depth: 2	0.676
combined predictions	0.715
min_samples_leaf: 2, with bagging	0.732

Introduction Variation from Random Features

```
# Fit a decision tree model to the "bag"
clf = DecisionTreeClassifier(random_state=1, min_samples_leaf=2,
                             splitter="random", max_features="auto")
clf.fit(bag[columns], bag["high_income"])
```

settings	test AUC
min_samples_leaf: 2	0.688
max_depth: 2	0.676
combined predictions	0.715
min_samples_leaf: 2, with bagging	0.732
min_samples_leaf: 2, with bagging and random subsets	0.734

Put it All Together

```
1 from sklearn.ensemble import RandomForestClassifier
2
3 clf = RandomForestClassifier(n_estimators=5, random_state=1,
4                             min_samples_leaf=2)
5 clf.fit(train[columns], train["high_income"])
6
7 predictions = clf.predict(test[columns])
8 print(roc_auc_score(test["high_income"], predictions))
```

0.7347461391939776

Reducing Overfitting

```
1 clf = DecisionTreeClassifier(random_state=1, min_samples_leaf=5)
2
3 clf.fit(train[columns], train["high_income"])
4
5 predictions = clf.predict(train[columns])
6 print(roc_auc_score(train["high_income"], predictions))
7
8 predictions = clf.predict(test[columns])
9 print(roc_auc_score(test["high_income"], predictions)).
```

0.8192570489534683

0.7139325899284541

Reducing Overfitting

```
1 clf = RandomForestClassifier(n_estimators=150, random_state=1,  
2                             min_samples_leaf=5)  
3 clf.fit(train[columns], train["high_income"])  
4  
5 predictions = clf.predict(train[columns])  
6 print(roc_auc_score(train["high_income"], predictions))  
7  
8 predictions = clf.predict(test[columns])  
9 print(roc_auc_score(test["high_income"], predictions))
```

0.7917047295143252

0.7498874343962398

Feature Importance

```
for score, name in sorted(zip(clf.feature_importances_, columns), reverse=True):  
    print('{} has a importance of {}'.format(name, score))
```



```
relationship has a importance of 0.35486065267354705  
education_num has a importance of 0.2238305467772321  
age has a importance of 0.16626910165605044  
hours_per_week has a importance of 0.09802596230288446  
occupation has a importance of 0.08434714139306755  
workclass has a importance of 0.039372047642714826  
race has a importance of 0.009447315730049248  
native_country has a importance of 0.009270711003041526  
sex has a importance of 0.007713429852542139  
marital_status has a importance of 0.006863090968870699
```

Case Study #2: predicting bike rentals

	instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	1	0	0	6	0	1	0.24	0.2879	0.81	0.0	3	13	16
1	2	2011-01-01	1	0	1	1	0	6	0	1	0.22	0.2727	0.80	0.0	8	32	40
2	3	2011-01-01	1	0	1	2	0	6	0	1	0.22	0.2727	0.80	0.0	5	27	32
3	4	2011-01-01	1	0	1	3	0	6	0	1	0.24	0.2879	0.75	0.0	3	10	13
4	5	2011-01-01	1	0	1	4	0	6	0	1	0.24	0.2879	0.75	0.0	0	1	1



EDA - Exploratory Data Analysis



Feature Engineering

```
def assign_label(hour):  
    if hour >=0 and hour < 6:  
        return 4  
    elif hour >=6 and hour < 12:  
        return 1  
    elif hour >= 12 and hour < 18:  
        return 2  
    elif hour >= 18 and hour <=24:  
        return 3
```

```
bike_rentals["time_label"] = bike_rentals["hr"].apply(assign_label)
```

Linear Regression vs Decision Tree vs Random Forest

```
from sklearn.linear_model import LinearRegression

predictors = list(train.columns)
predictors.remove("cnt")
predictors.remove("casual")
predictors.remove("registered")
predictors.remove("dteday")

reg = LinearRegression()

reg.fit(train[predictors], train["cnt"])

import numpy
predictions = reg.predict(test[predictors])

np.sqrt(np.mean((predictions - test["cnt"]) ** 2))
```

129.68788758633

Linear Regression vs Decision Tree vs Random Forest

```
from sklearn.tree import DecisionTreeRegressor  
reg = DecisionTreeRegressor(min_samples_leaf=2)  
reg.fit(train[predictors], train["cnt"])  
predictions = reg.predict(test[predictors])  
np.sqrt(np.mean((predictions - test["cnt"]) ** 2))
```

51.88006941722633

Linear Regression vs Decision Tree vs Random Forest

```
from sklearn.ensemble import RandomForestRegressor

reg = RandomForestRegressor(n_estimators=100,min_samples_leaf=2)
reg.fit(train[predictors], train["cnt"])

predictions = reg.predict(test[predictors])

np.sqrt(np.mean((predictions - test["cnt"]) ** 2))
```

38.784474743954746

When to use Random Forest

- Strengths of a Random Forest
 - Very accurate predictions
 - Resistance to overfitting
- Weakness
 - They are difficult to interpret
 - They take longer to create

Lesson #22 - Random Forest.ipynb

